

Deep Learning

Data mining: Advanced Aspects

Siham Tabik
University of Granada
siham@ugr.es

Schedule

Deep Learning		
February 13 th	February 14 th	February 20 th
Deep Neural Networks. <ul style="list-style-type: none">• Convolutional Neural Networks (CNNs)• Classification & detection• Two case studies	Tensorflow. <ul style="list-style-type: none">• Get familiar• Build our first CNN model	Recurrent Neural Networks. <ul style="list-style-type: none">• Some theory• an example using tensorflow• One case study

Logistics

- Forum for debates, ...
https://github.com/SihamTabik/Patio_DL/issues
- Textbook “Deep Learning with Tensorflow”. Please, download the pdf from PRADO.

Convolutional Neural Networks

Siham Tabik
University of Granada
siham@ugr.es

Outline

1. What are ANN and CNNs?
2. How do CNNs work?
 - Convolution layers, pooling layers, FC layers, gradient, Backpropagation
3. Transfer learning
4. Data augmentation
5. One case study of CNN-classification model
6. CNNs for Detection
7. Two case studies

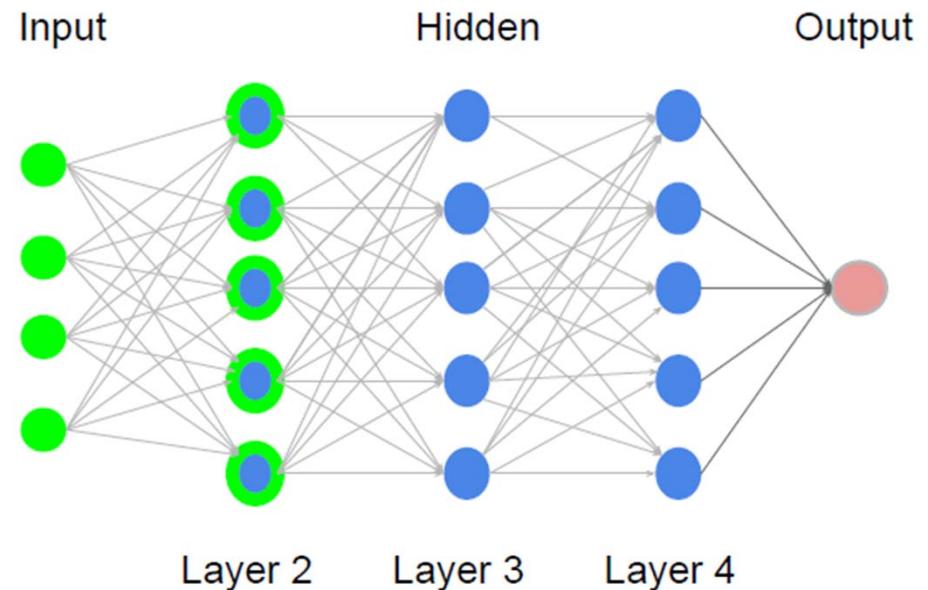
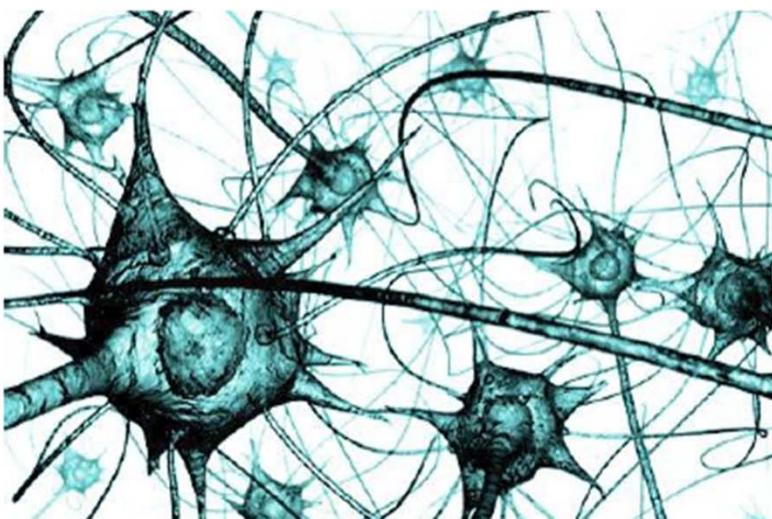
Outline

1. What are ANN and CNNs?
2. How do CNNs work?
 - Convolution layers, pooling layers, FC layers, gradient, Backpropagation
3. Transfer learning
4. Data augmentation
5. One case study of CNN-classification model
6. CNNs for Detection
7. Two case studies

Artificial Neural Networks

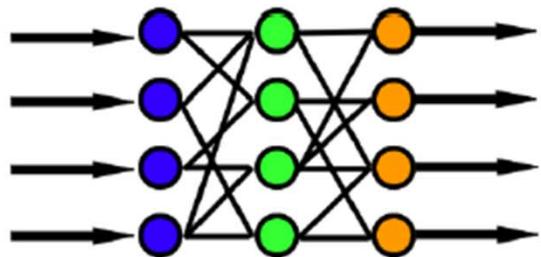
Machine learning algorithms

Learn and predict on data

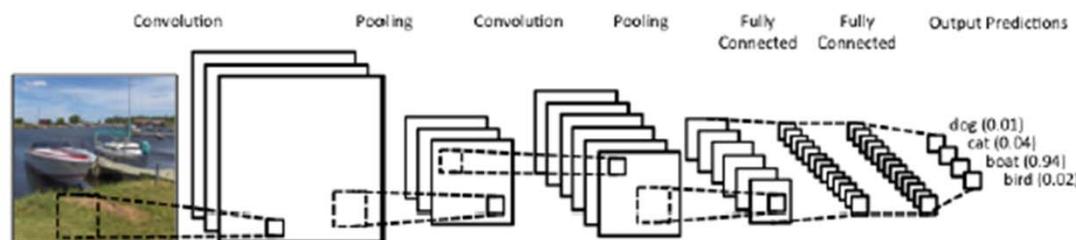


Artificial Neural Networks

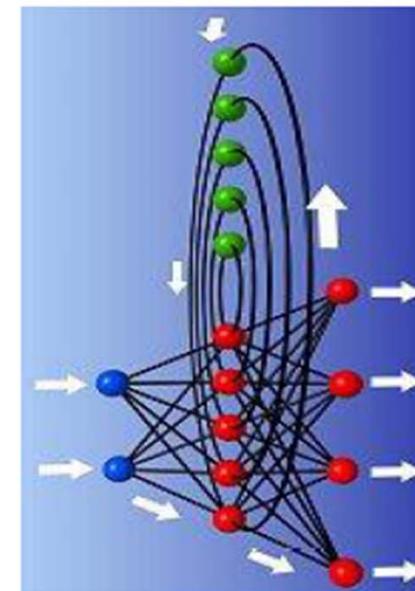
Feed Forward Nets (FNN)



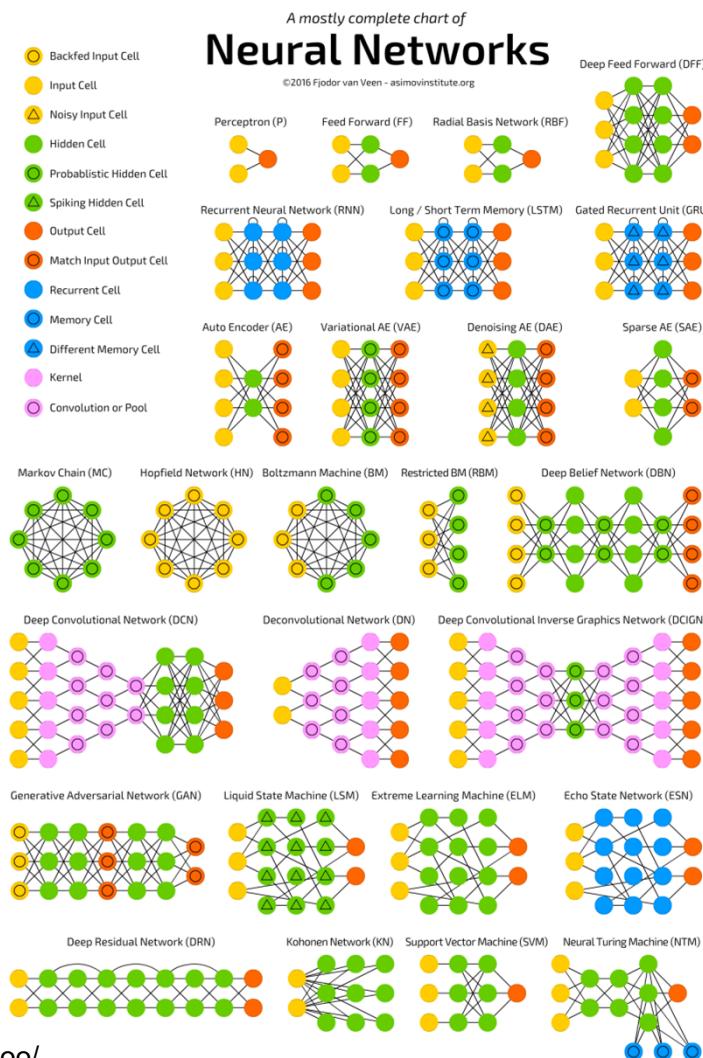
Convolutional Nets (CNN)



Recurrent Nets (RNN)



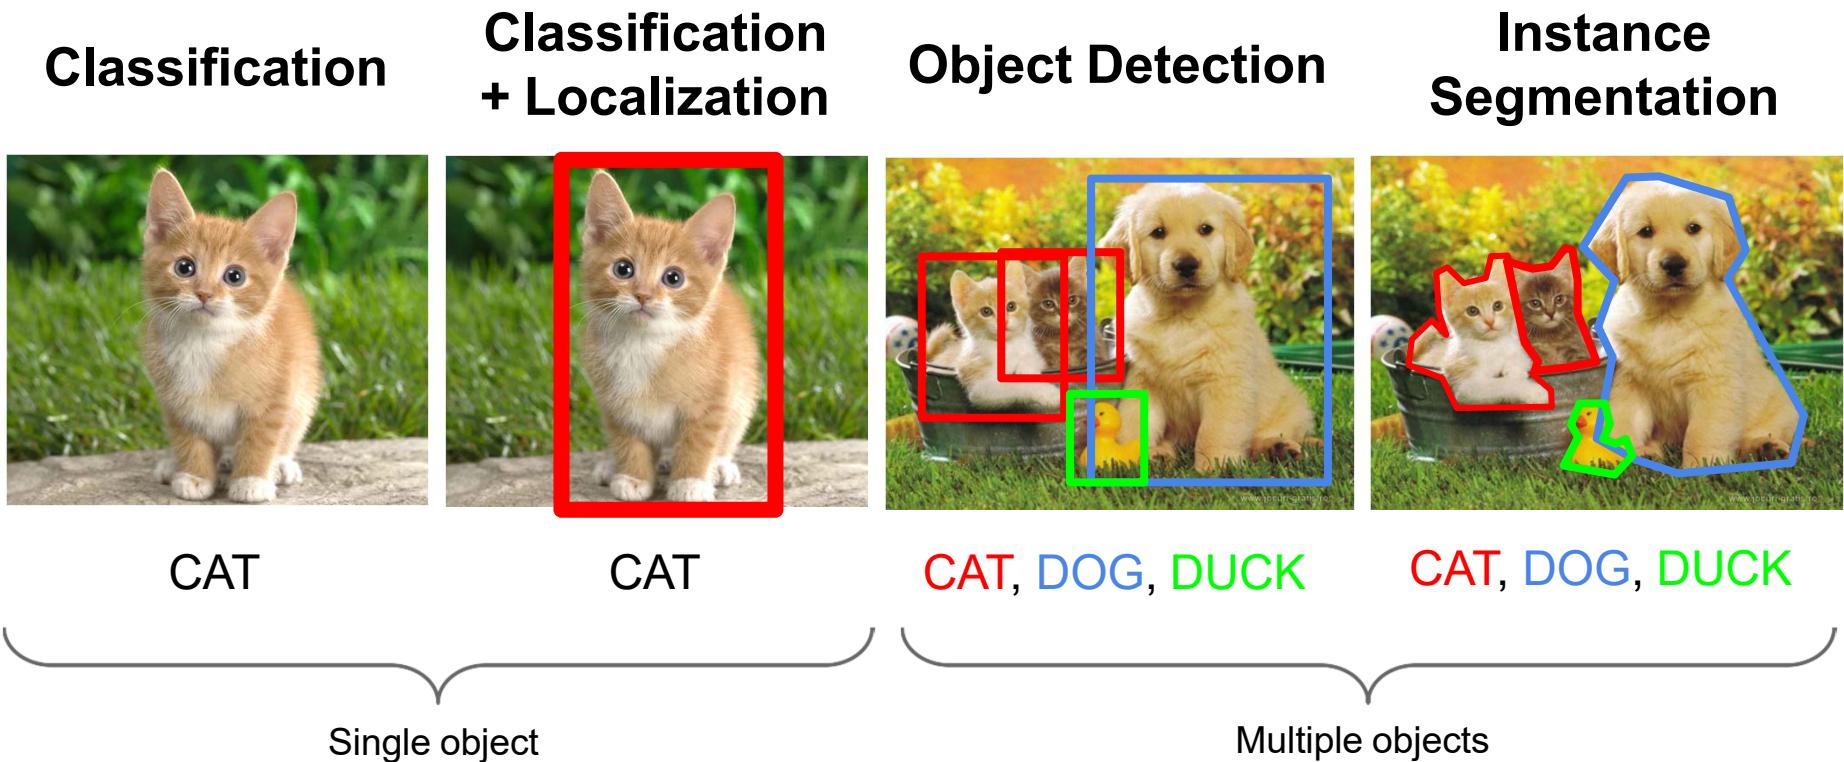
... and others



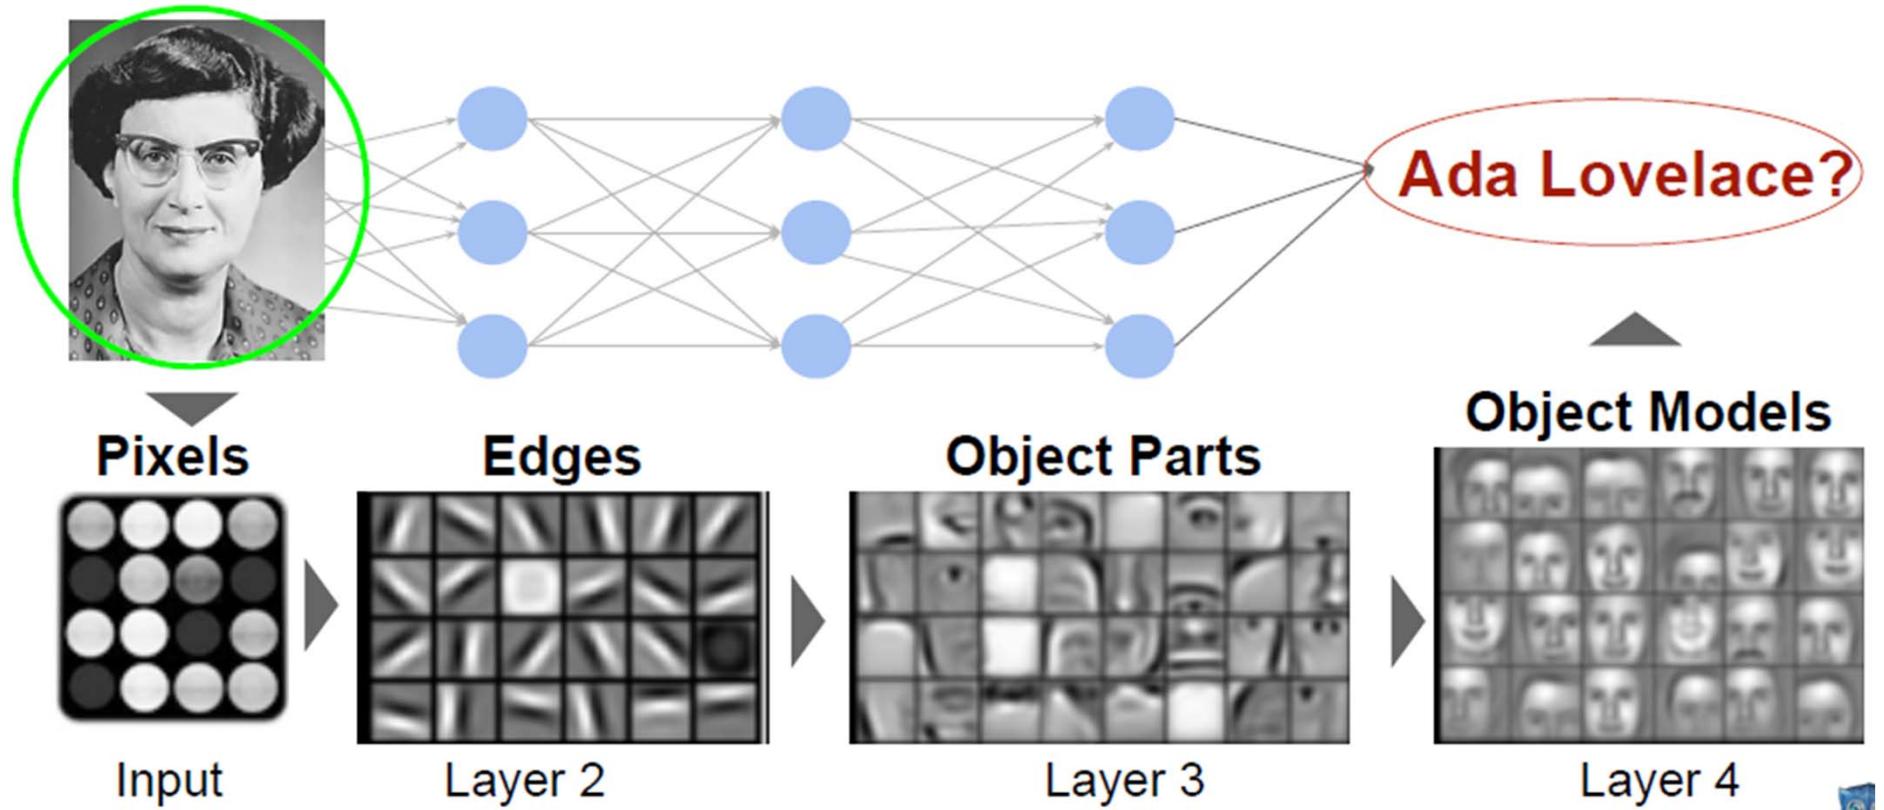
Source:

<http://www.asimovinstitute.org/neural-network-zoo/>

By the way, what is image classification?



Convolutional Neural Networks



PASCAL Visual Object Classes Challenge (2005-2012)



Aeroplanes



Bicycles



Birds



Boats



Bottles



Bus



Cars



Cat



Chair



Cow

- **Database:** Public dataset of 10,103 images & 20 object classes
- **Annual Competition:** PASCAL VOC Challenge
- **Networks:** The most accurate net wins the challenge



Dining tables



Dogs



Horses



Motorbikes



People



Potted plants



Sheep



Sofas



Trains



TV Monitors

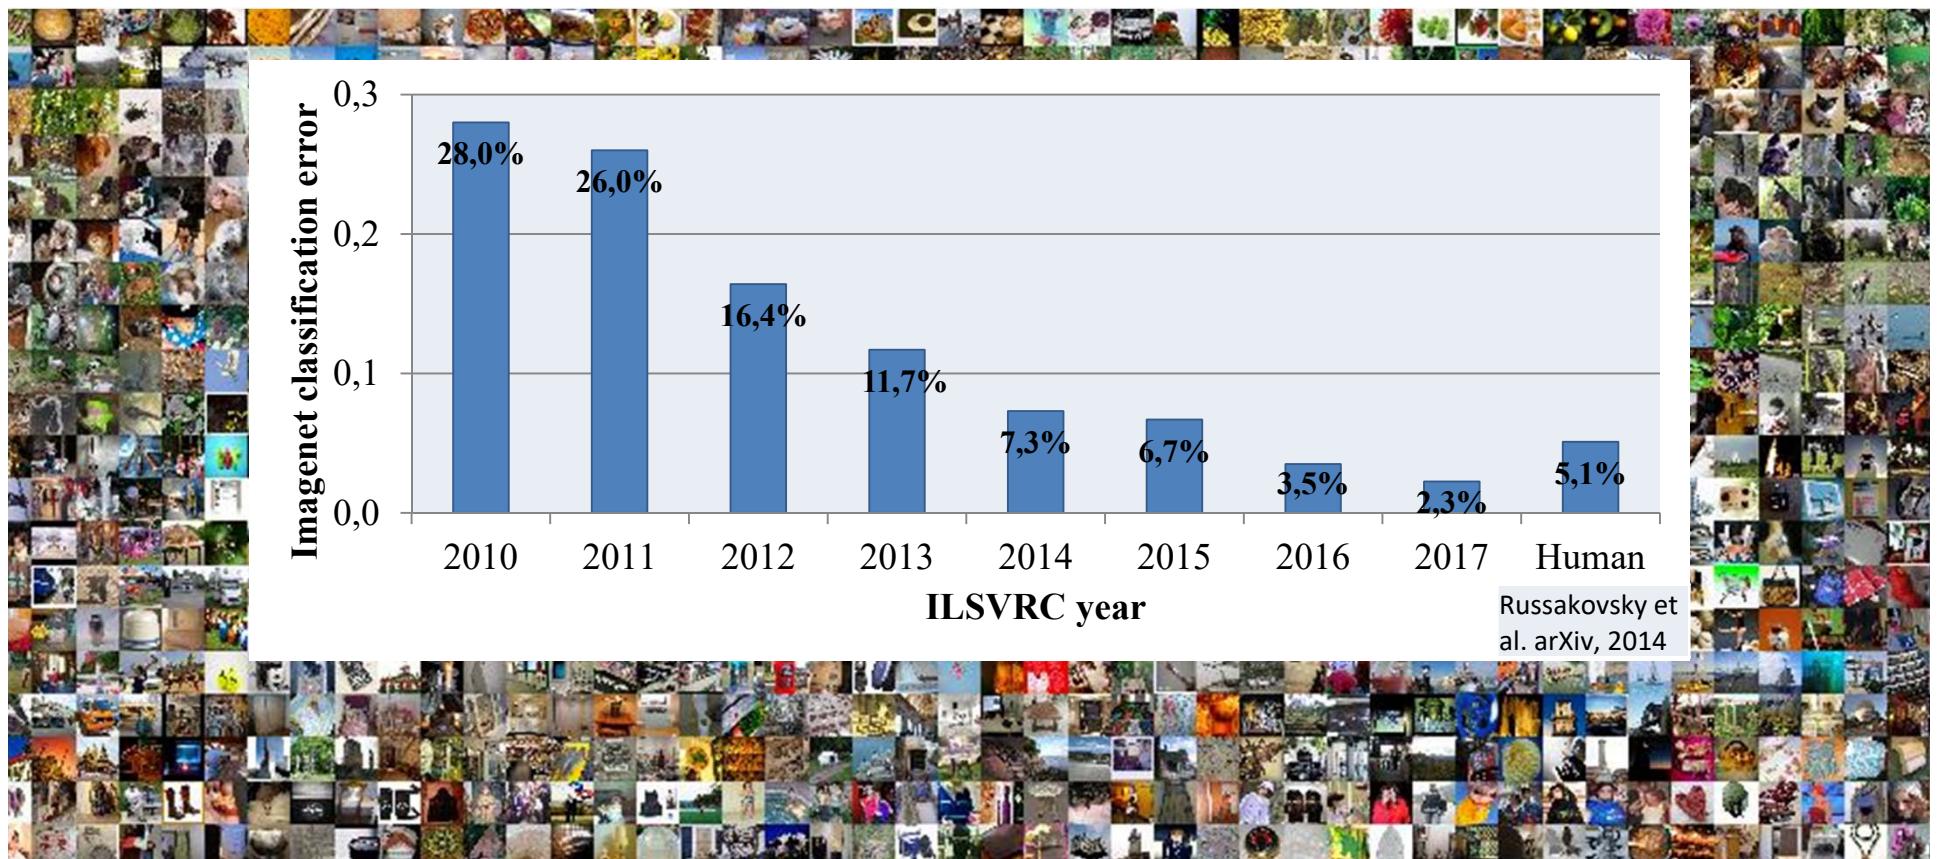
IMAGENET image classification challenge



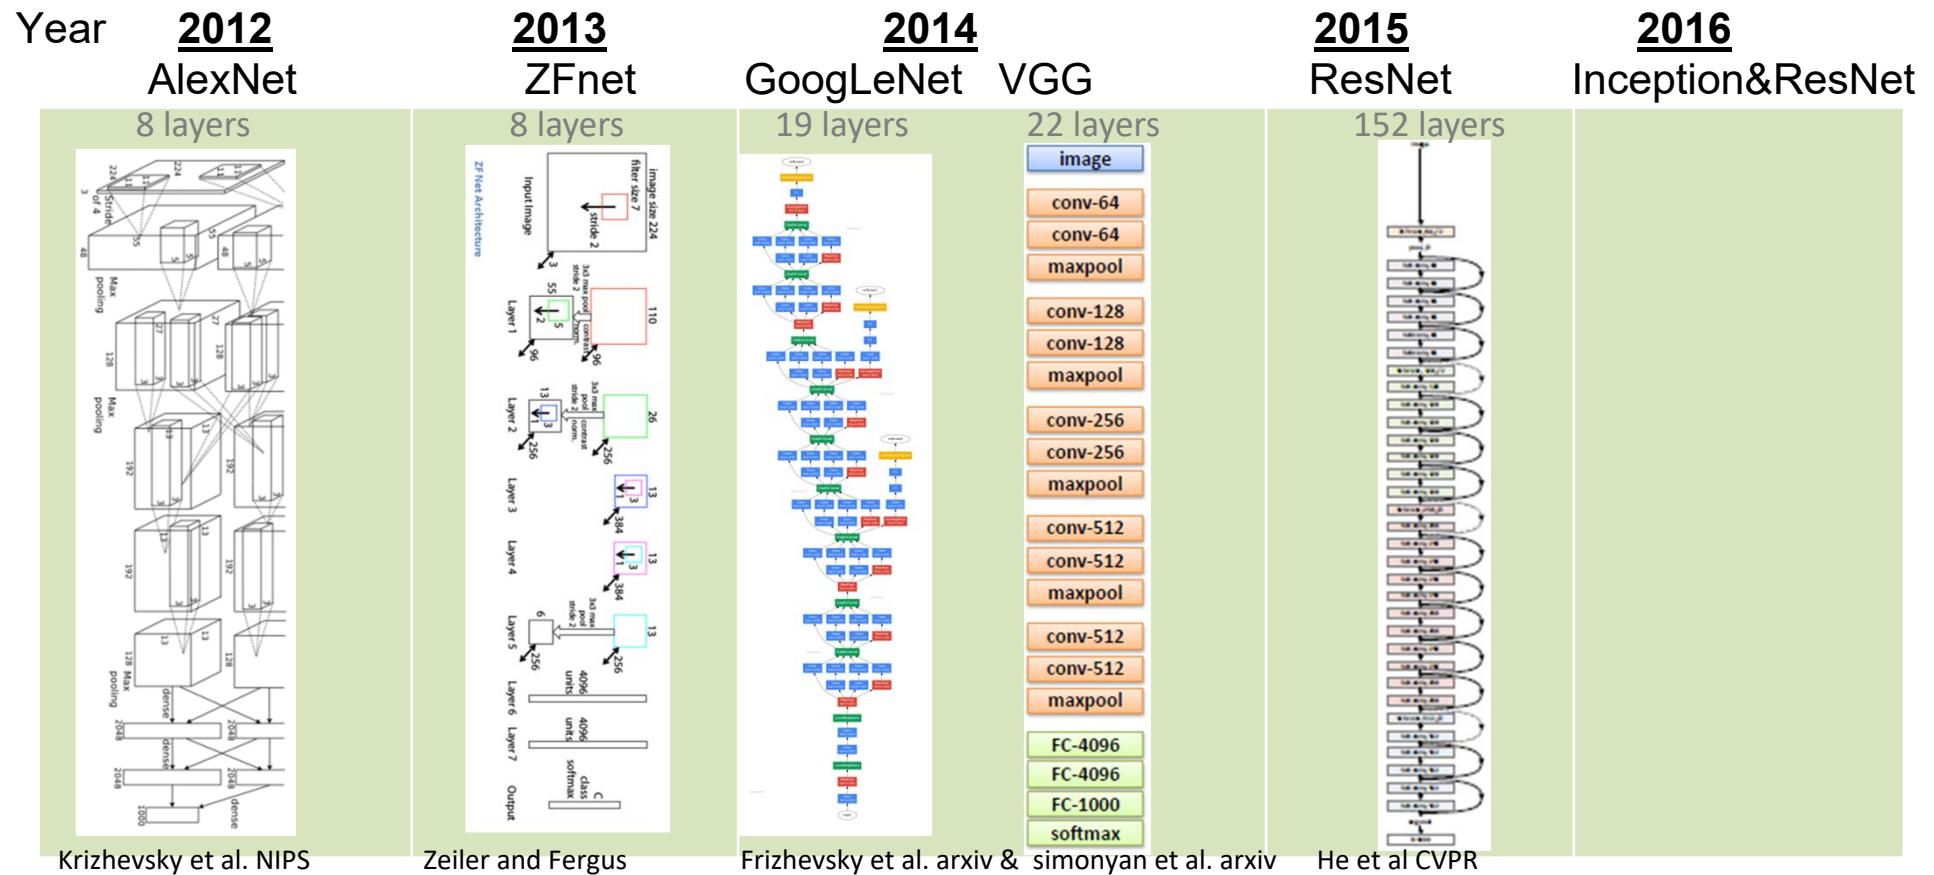
- **Database:** Public dataset of 14,197,122 images & 21,841 classes
- **Annual Competition:** The Large Scale Visual Recognition Challenge (ILSVRC)
 - **Classification task:** 1,431,167 images & 1000 classes (1.2M train+100.000 test)
 - **Detection task, segmentation task, ...**
- **Networks:** The most accurate net win the challenge



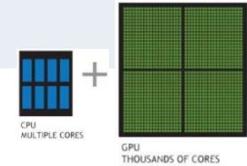
IMAGENET image classification challenge



IMAGENET image classification challenge



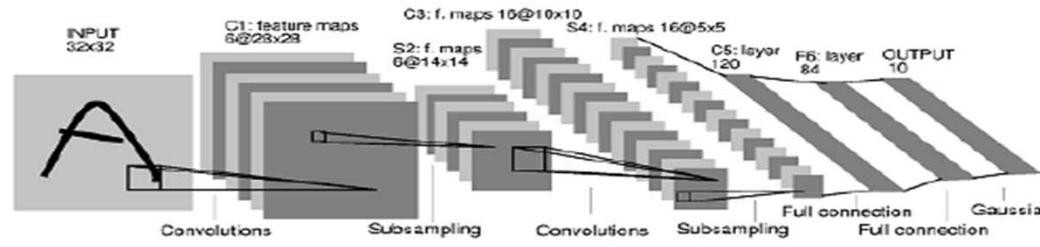
Evolution of CNNs

	1975	1985	1998	2009	2012
CNNs	1 st implementation of Conv. & pooling layers by K. Fukushima	1 st CNN trained with Backpropagation	1 st CNN trained with backp. on MNIST: LeNet (0.8% error) by Y. LeCun		Won the ILSVRC classification challenge
Learning Algorithms	Costly	Backpropagation	Backpropagation	Backpropagation	Backpropagation
Database	Small	Small	MNIST	1 st massive dataset ImageNet 	
Hardware	Limited capacity	Limited capacity	Limited capacity	Powerful GPUs and CPUs 	

CNNs architecture

1998

LeCun et al.

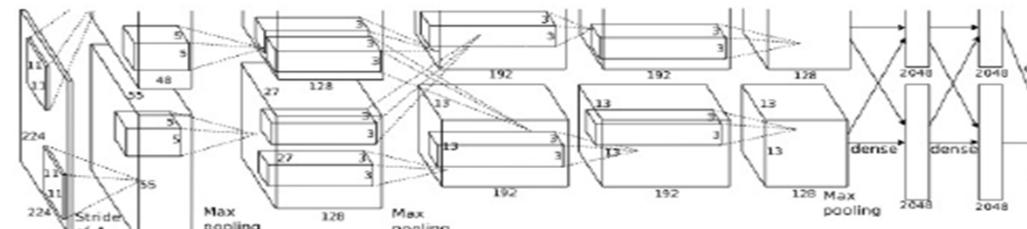


of transistors
 10^6
pentium II

of pixels used in training
 10^7 NIST

2012

Krizhevsky et al.



of transistors
 10^9



of pixels used in training
 10^{14} IMAGENET

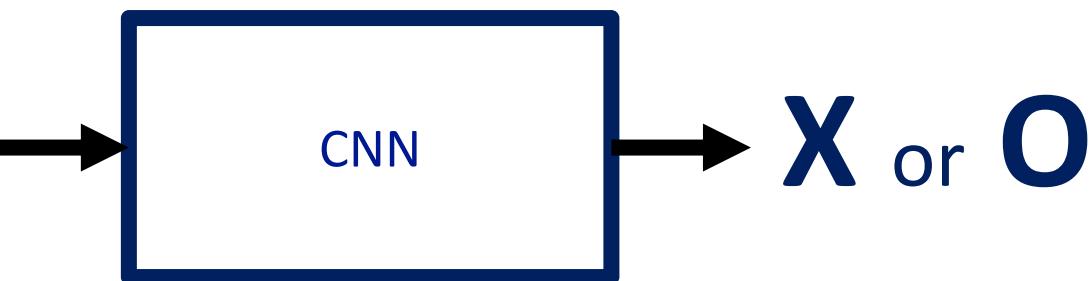
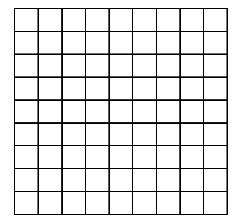
Outline

1. What are ANN?
2. What are CNNs?
3. How CNNs work?
 - Convolution layers, pooling layers, FC layers, Gradient, Backpropagation
4. Transfer learning
5. Data augmentation
6. One case study
7. CNNs for Detection
8. Two case studies

How CNNs work?

A toy CNN: says whether a picture is of an X or an O

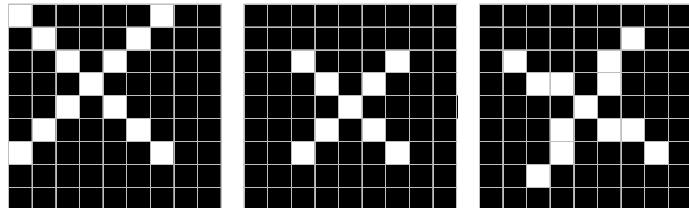
A two-dimensional
array of pixels



For example



Trickier cases

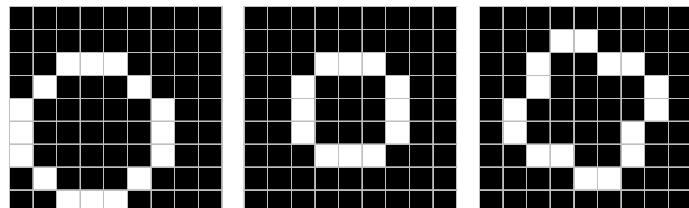


translation

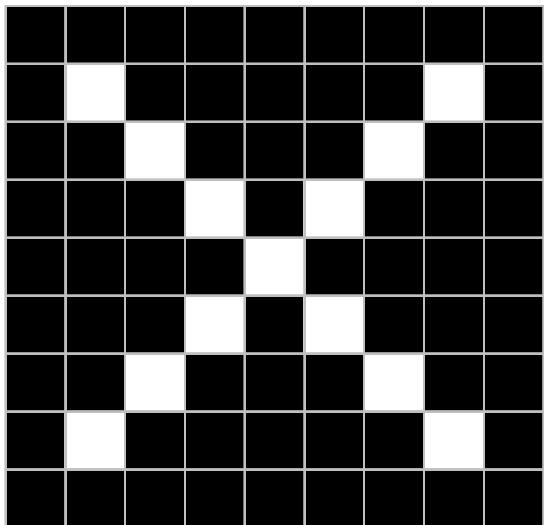
scaling

rotation

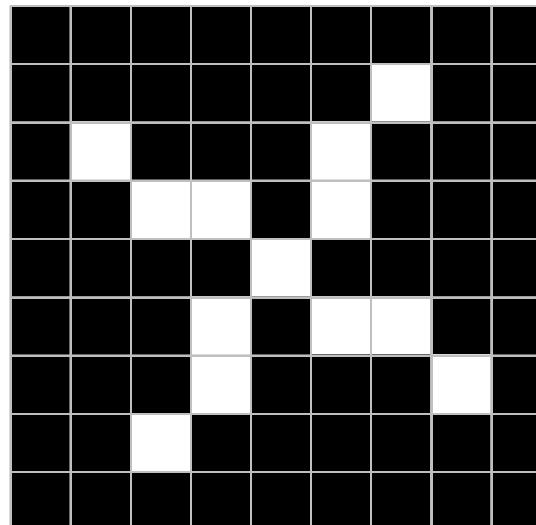
weight



Deciding is hard



?



Deciding is hard

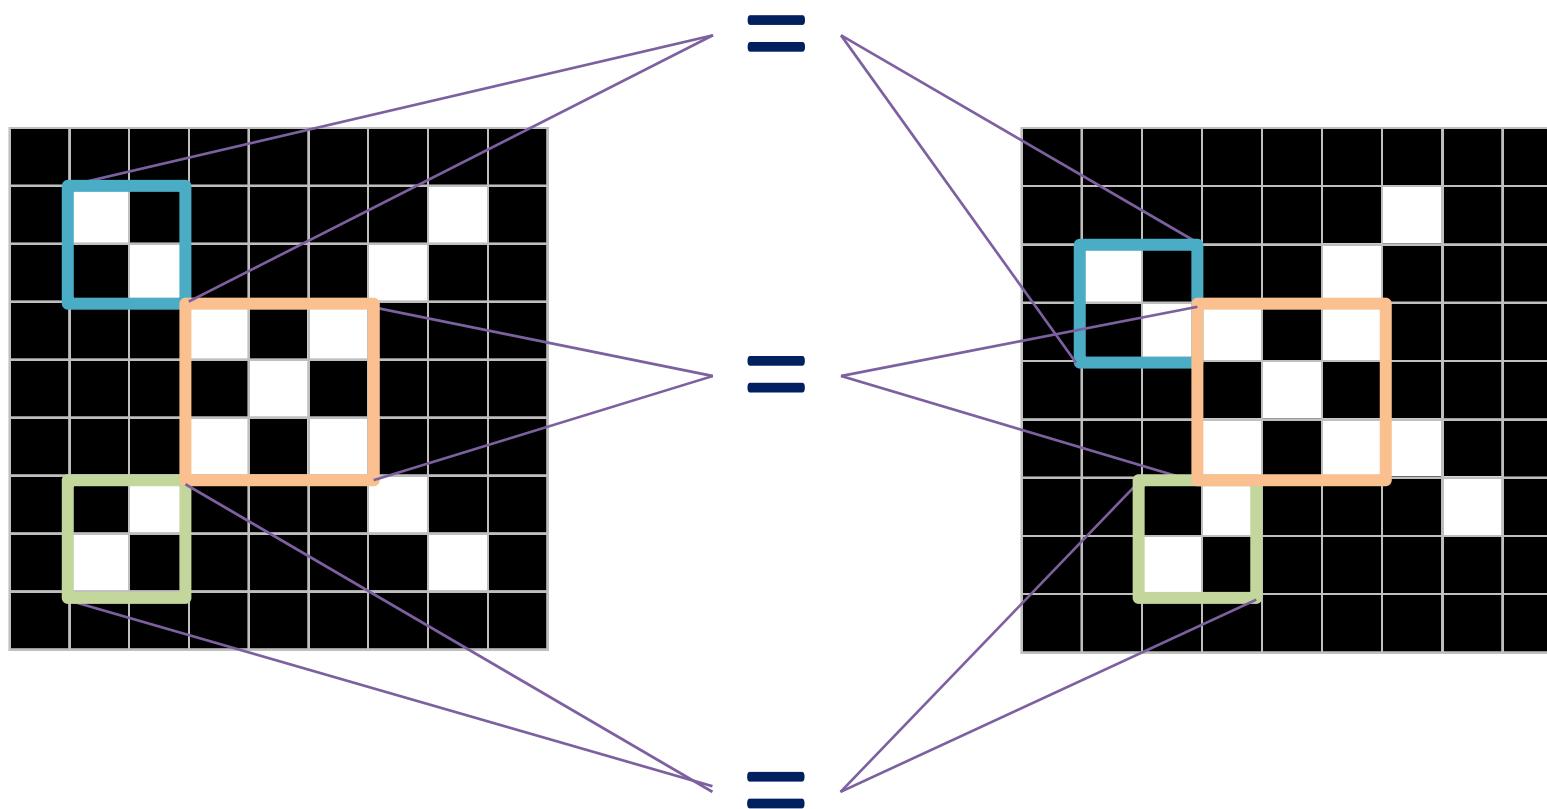


What computers see

Computers are literal



CNNs match pieces of the image



Features match pieces of the image

1	-1	-1
-1	1	-1
-1	-1	1

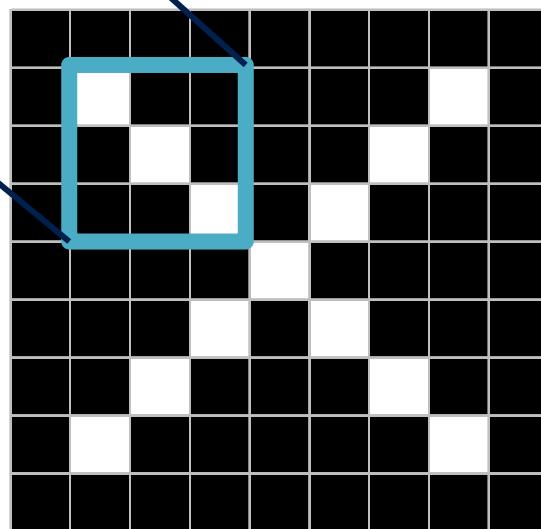
1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

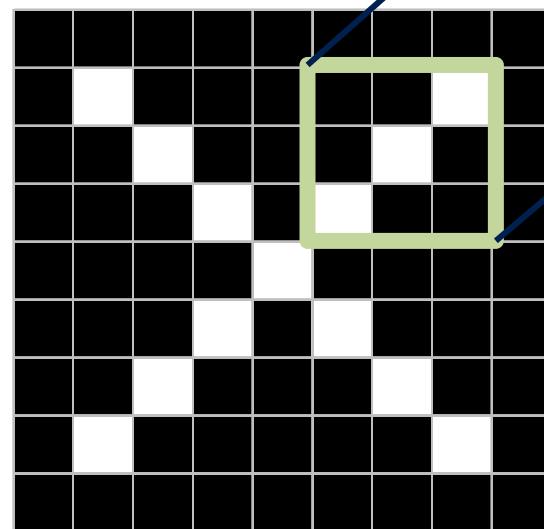
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

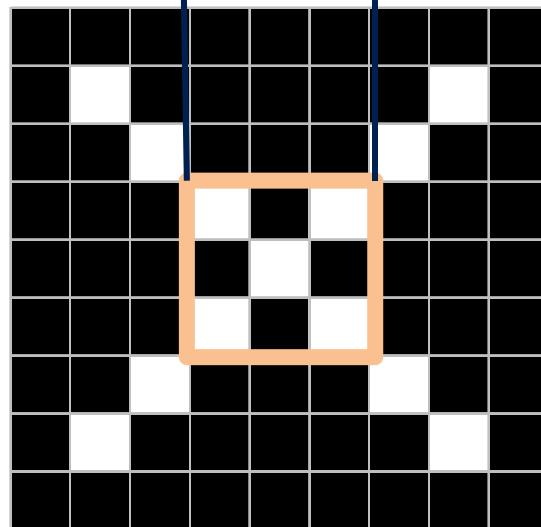
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

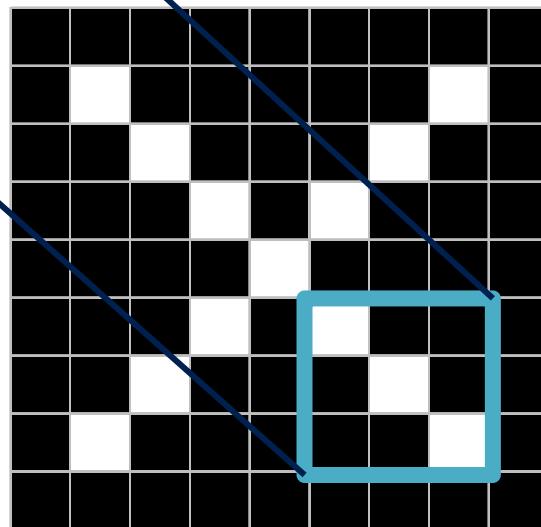
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

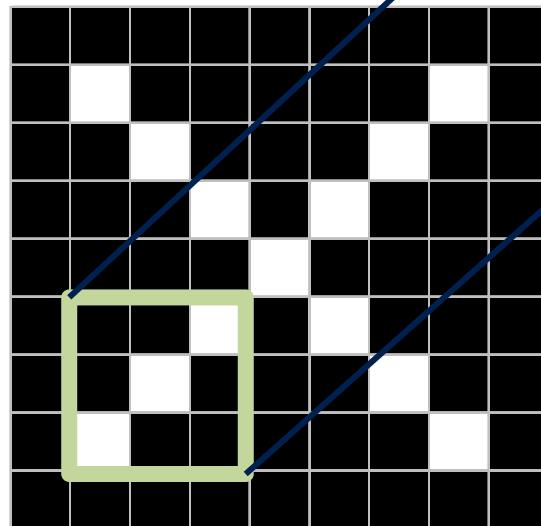
-1	-1	1
-1	1	-1
1	-1	-1



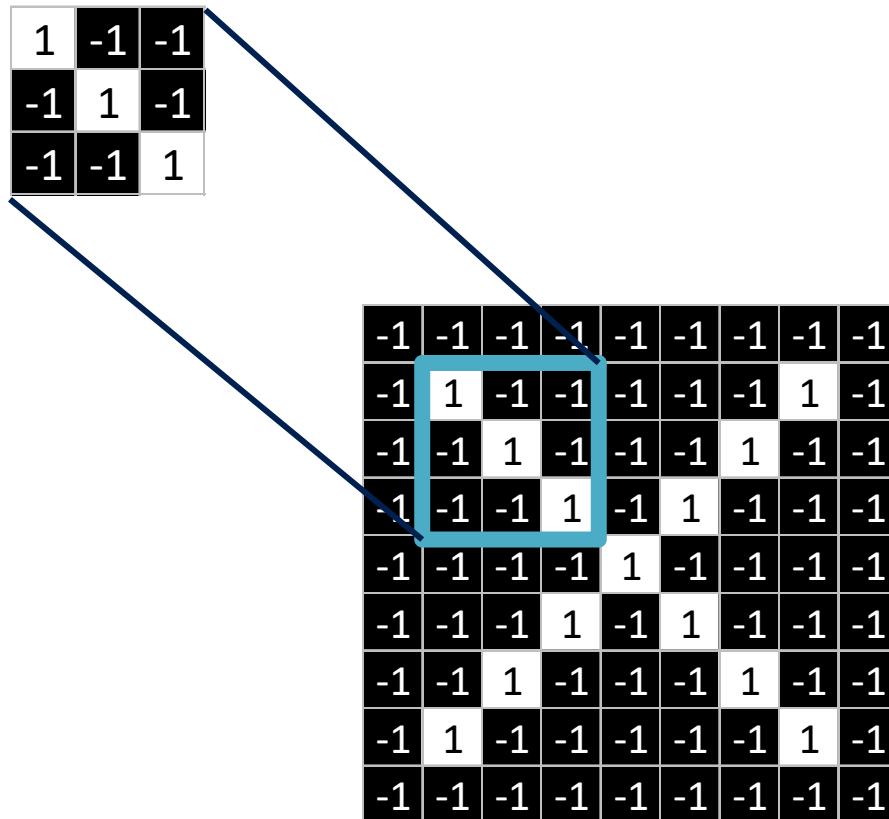
1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1



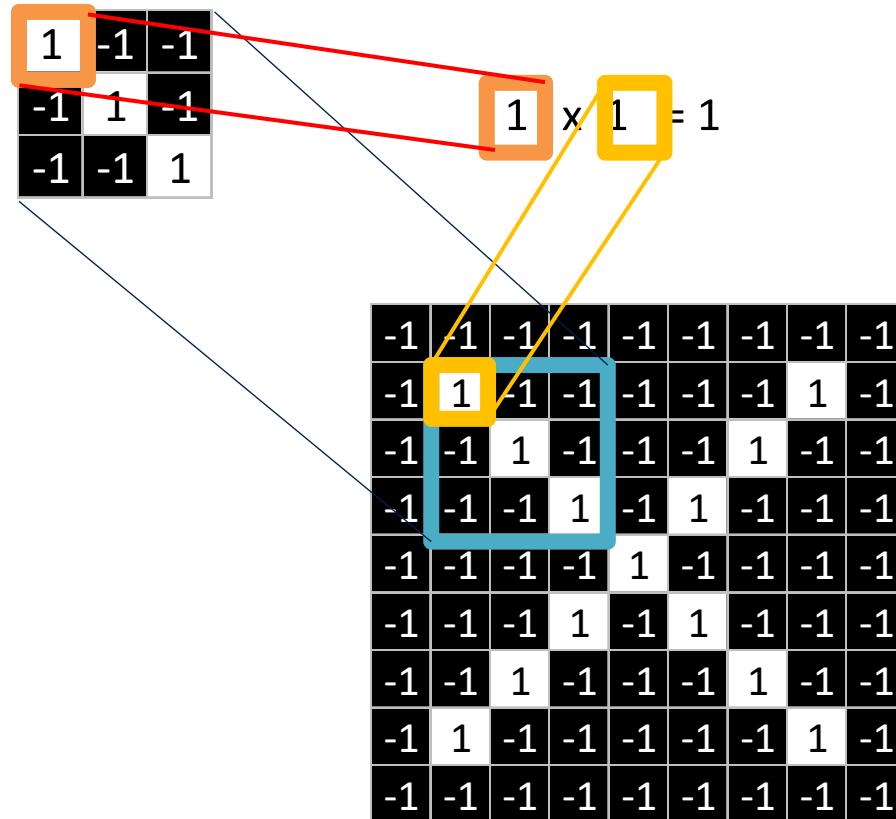
Filtering: The math behind the match



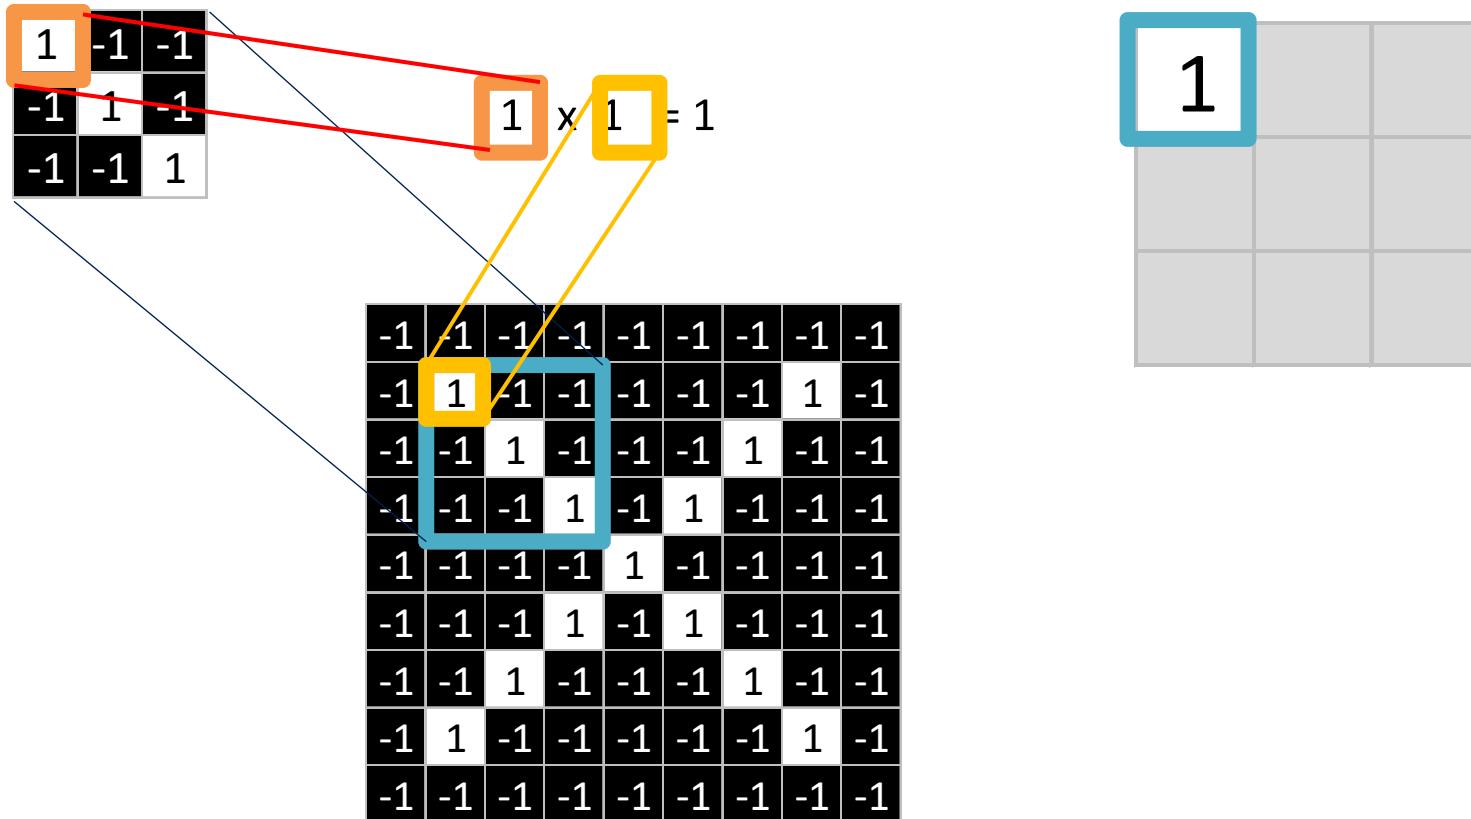
Filtering: The math behind the match

1. Line up the feature and the image patch.
2. Multiply each image pixel by the corresponding feature pixel.
3. Add them up.
4. Divide by the total number of pixels in the feature.

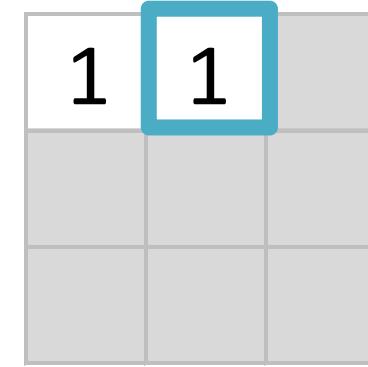
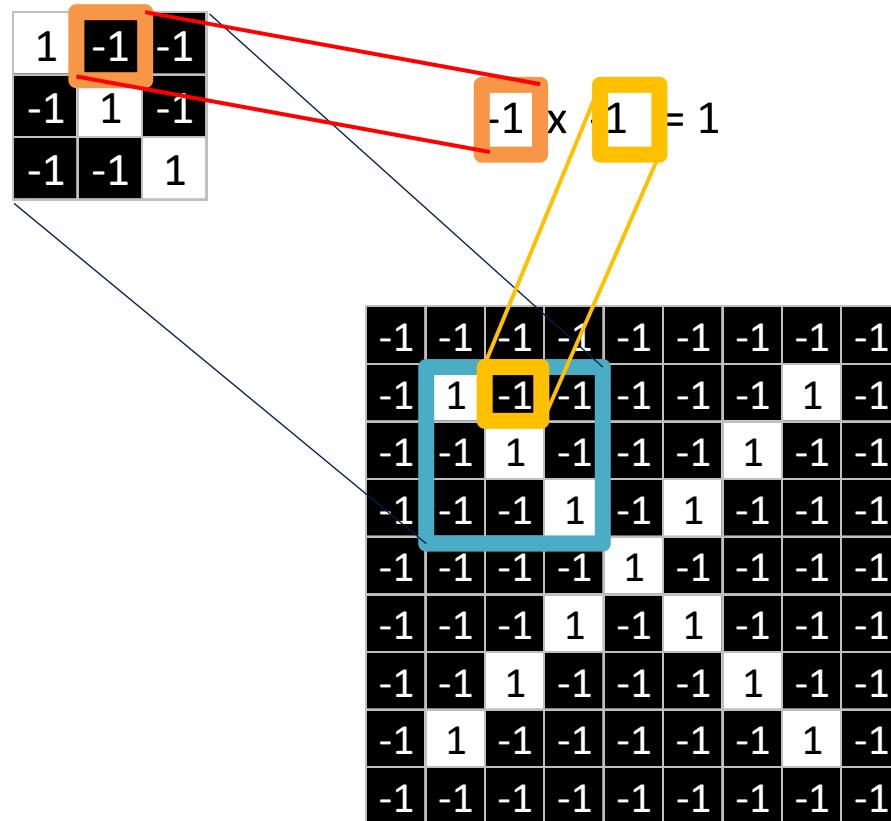
Filtering: The math behind the match



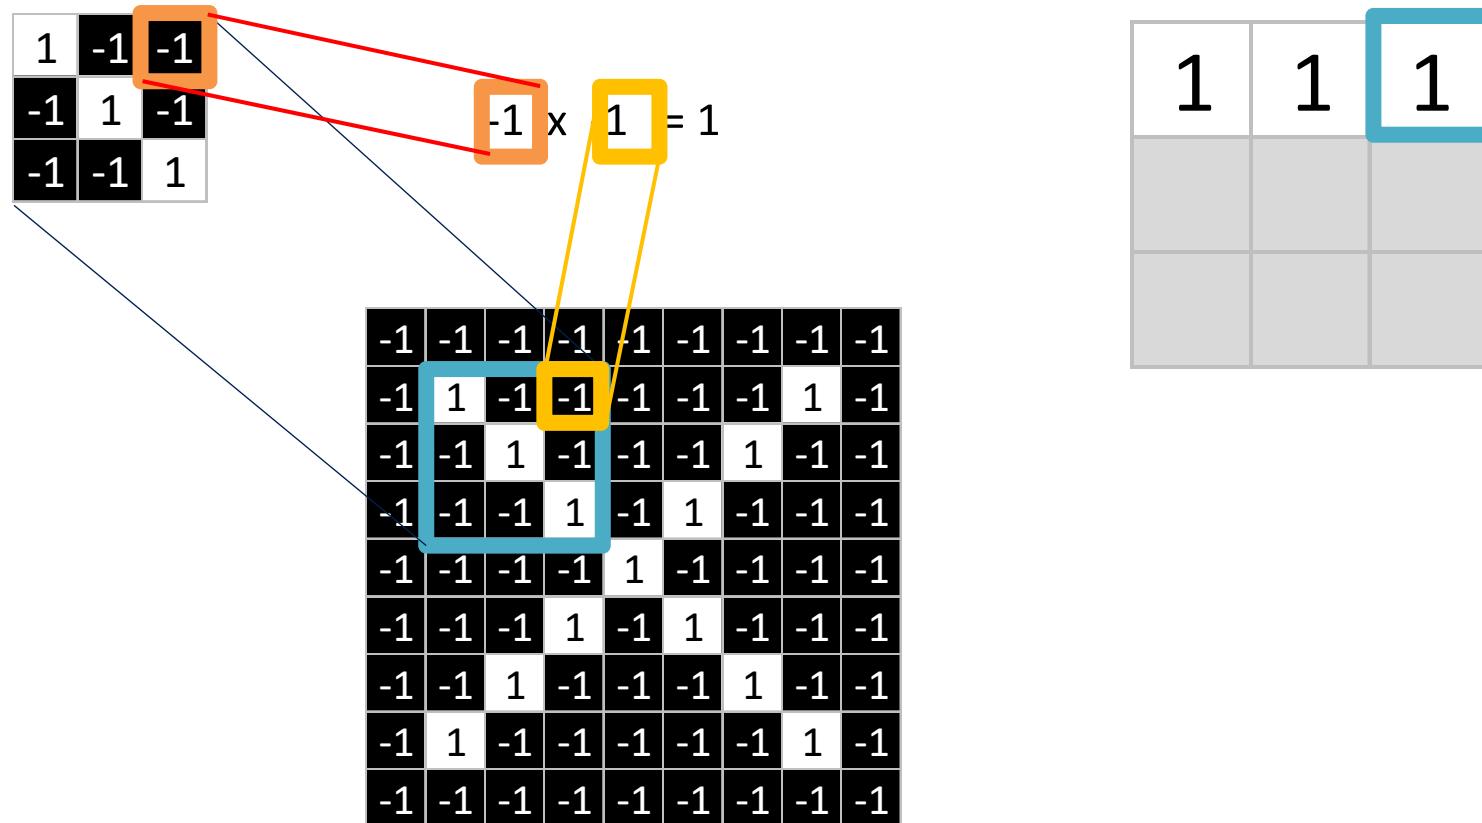
Filtering: The math behind the match



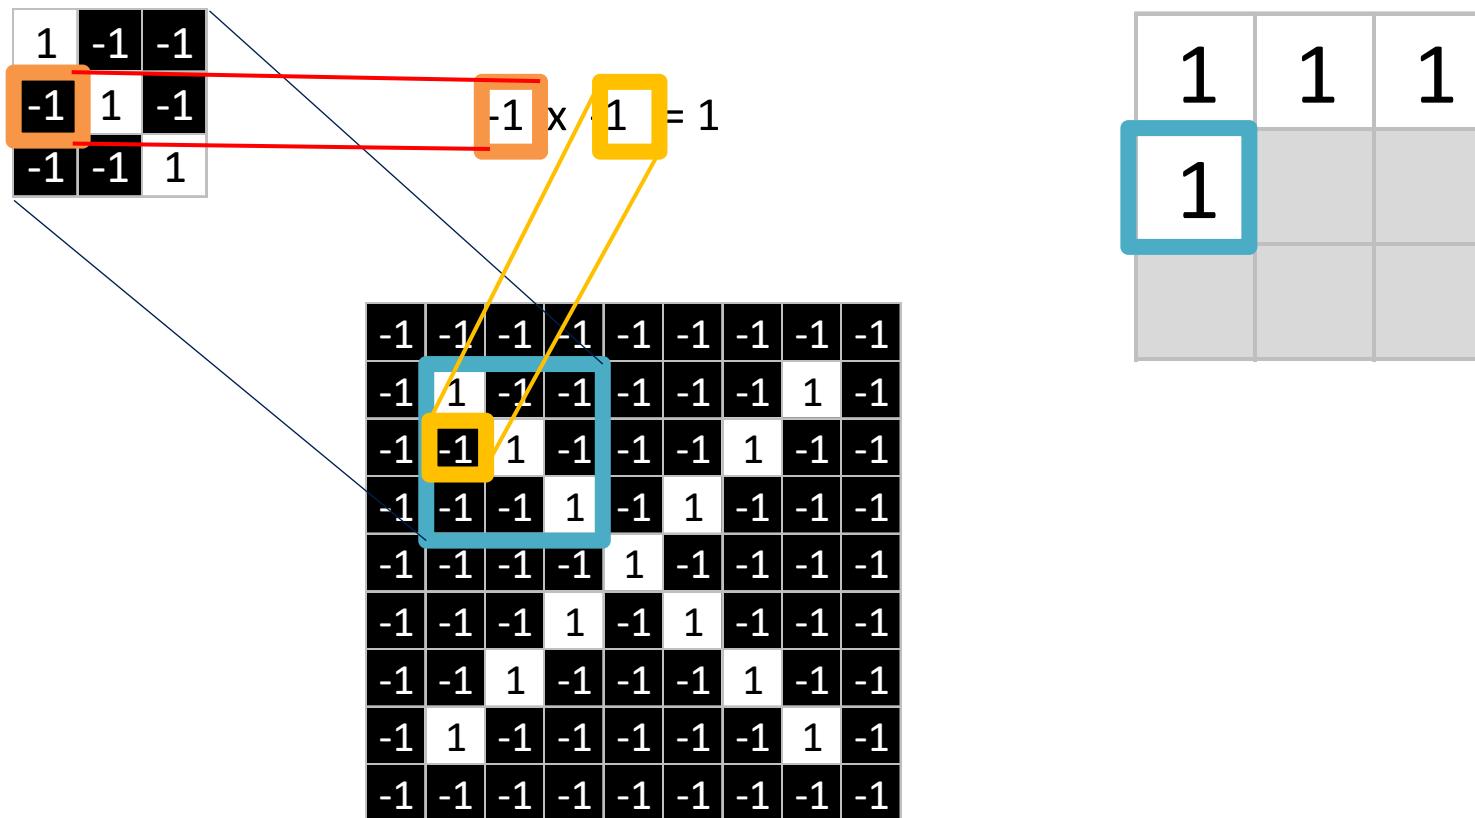
Filtering: The math behind the match



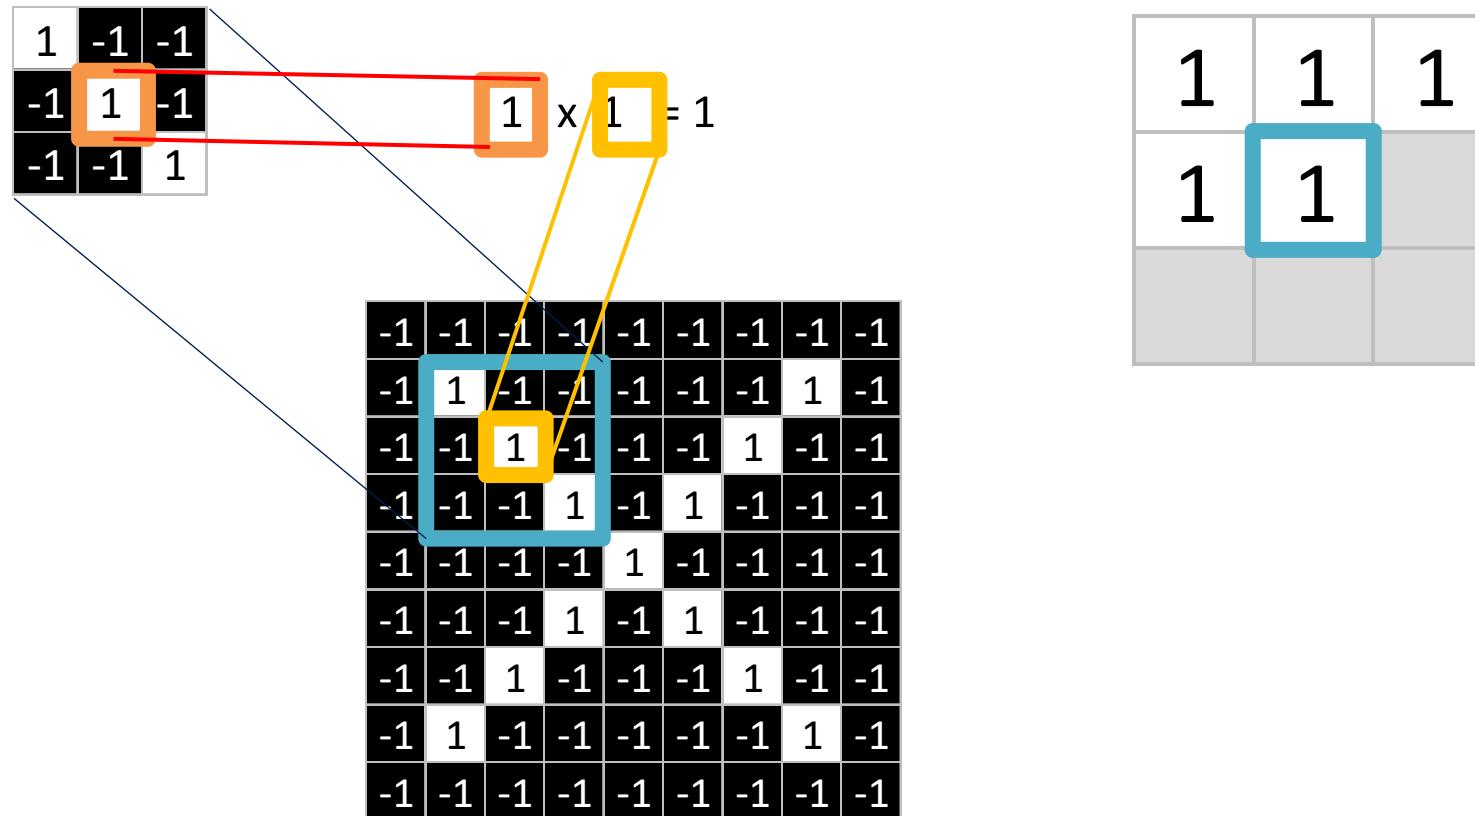
Filtering: The math behind the match



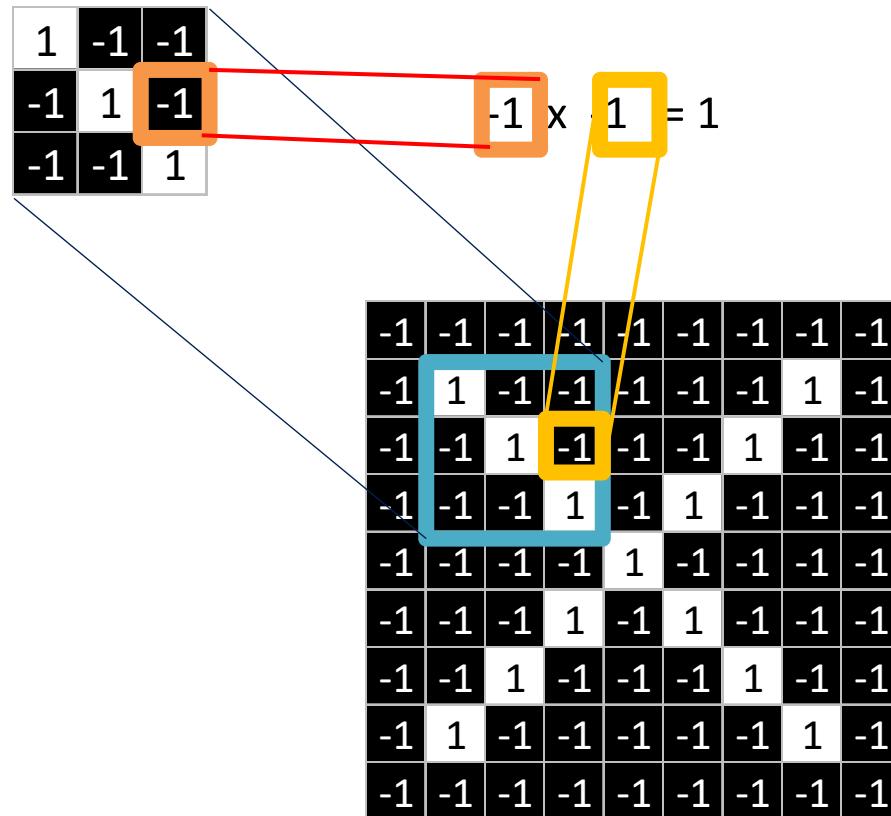
Filtering: The math behind the match



Filtering: The math behind the match

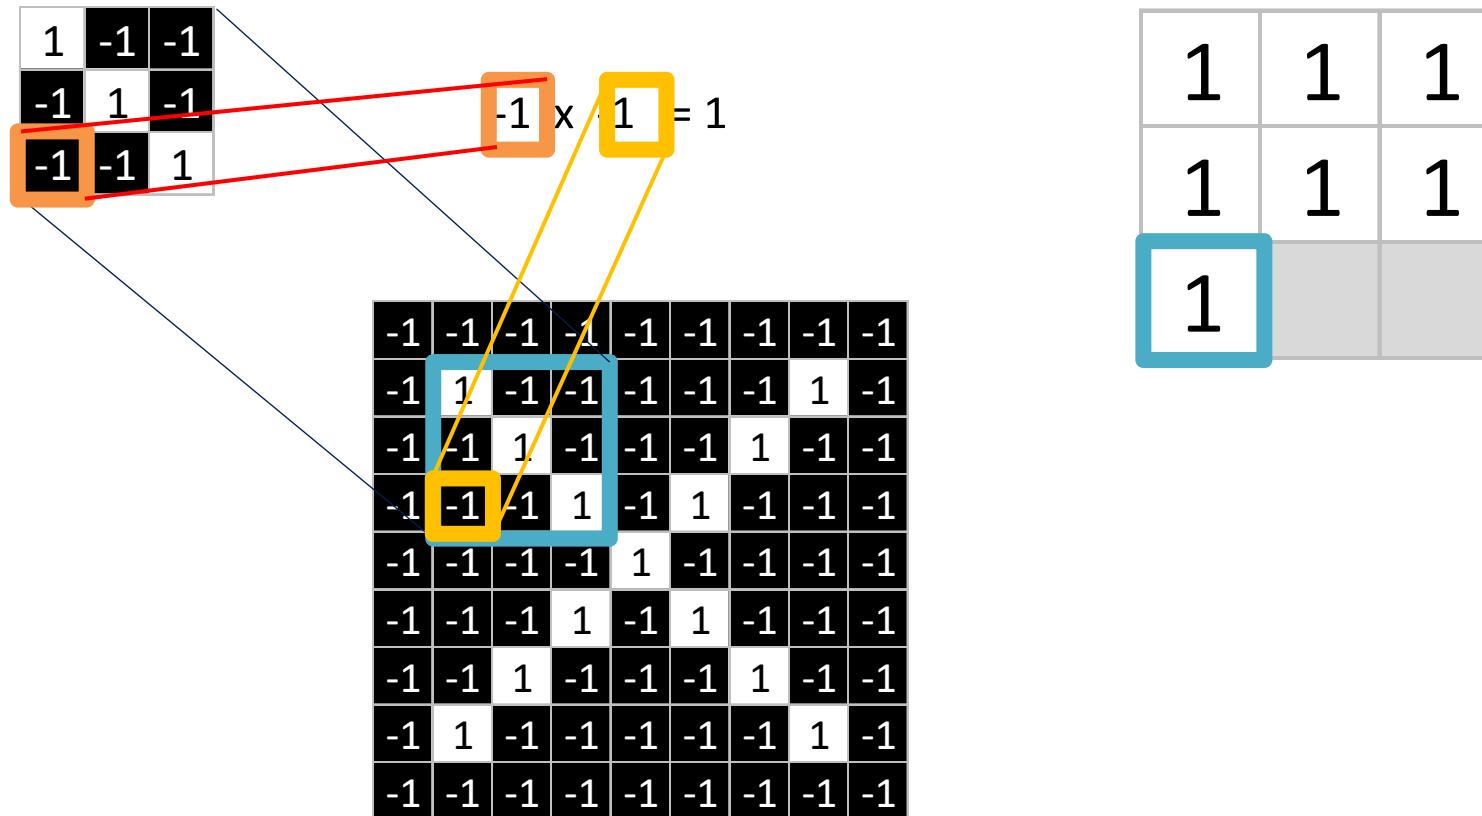


Filtering: The math behind the match

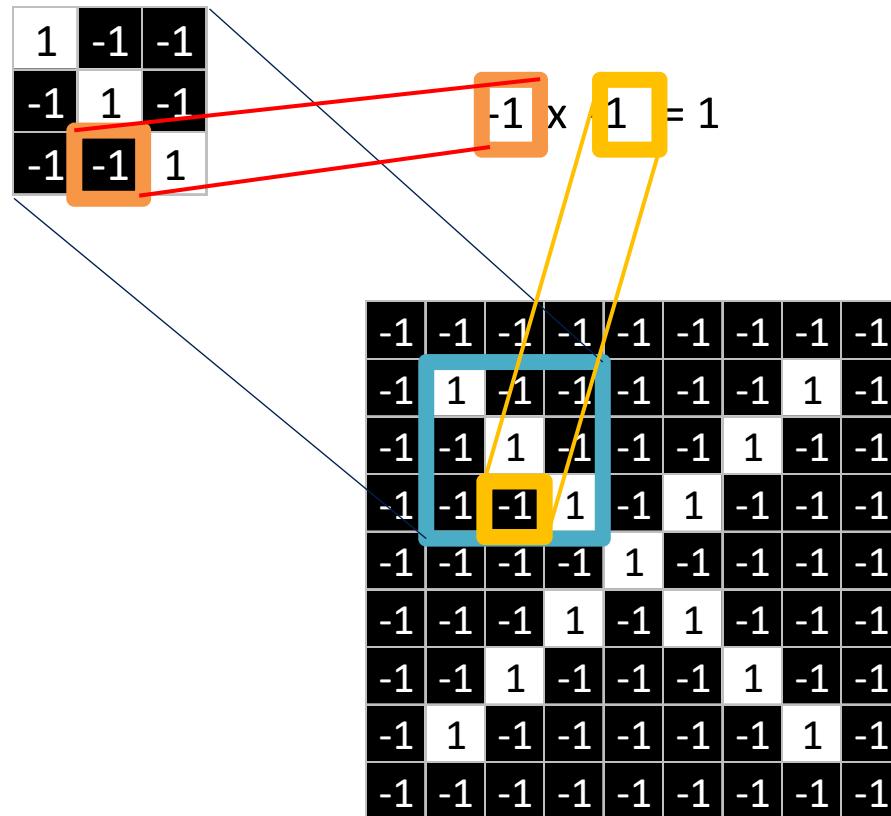


1	1	1
1	1	1
1	1	1

Filtering: The math behind the match

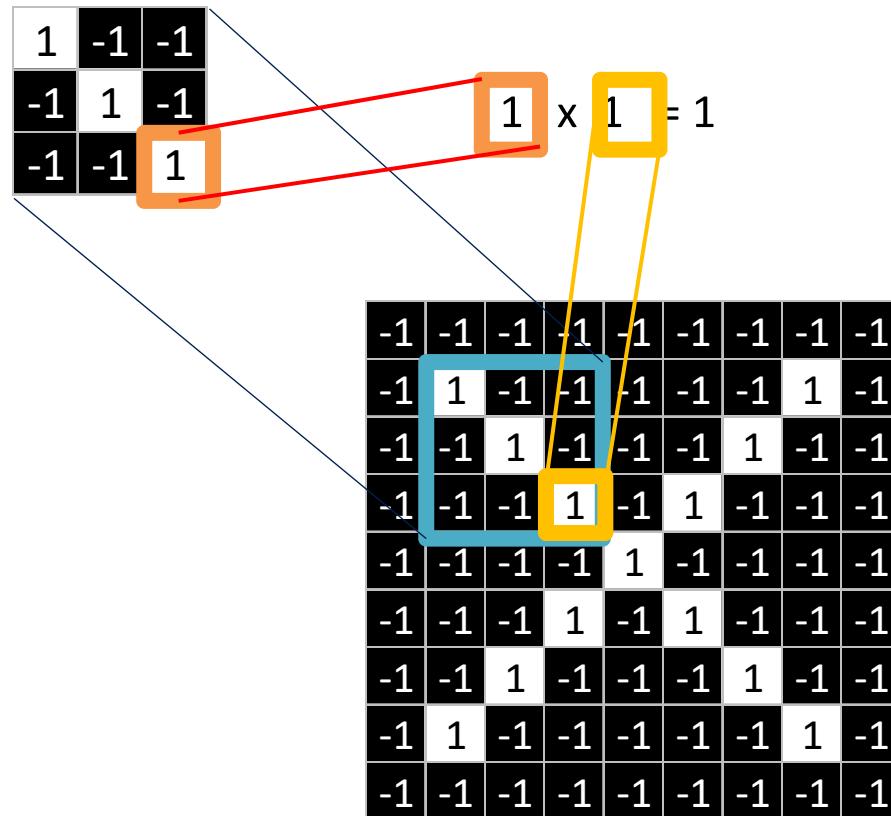


Filtering: The math behind the match



1	1	1
1	1	1
1	1	1

Filtering: The math behind the match



1	1	1
1	1	1
1	1	1

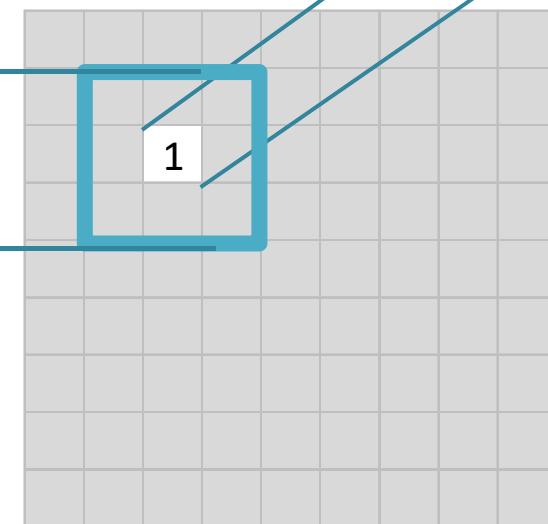
Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

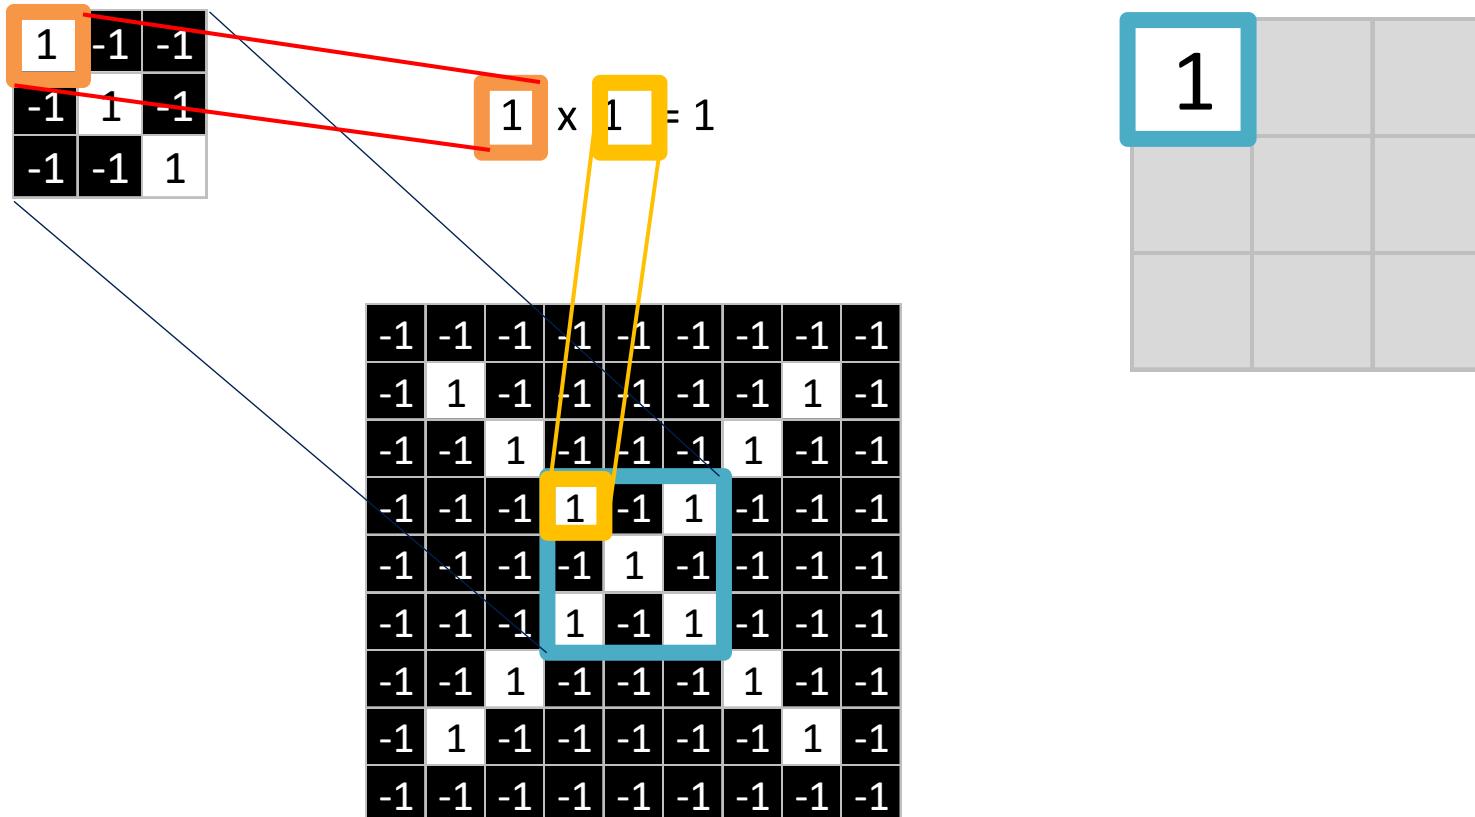
1	1	1
1	1	1
1	1	1

$$\begin{array}{r} 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 \\ \hline 9 \end{array} = 1$$

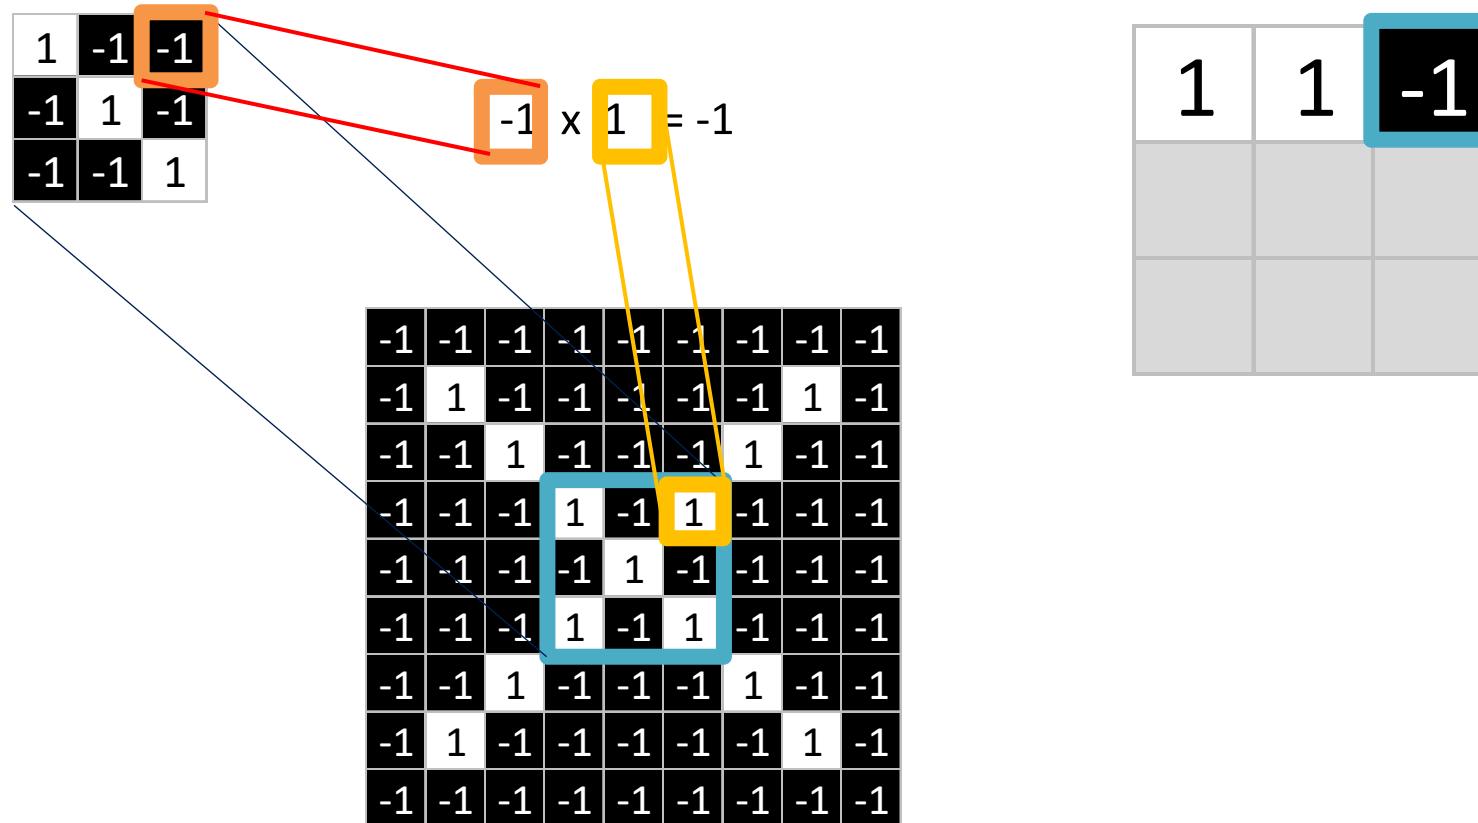
-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



Filtering: The math behind the match



Filtering: The math behind the match



Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	-1
1	1	1
-1	1	1

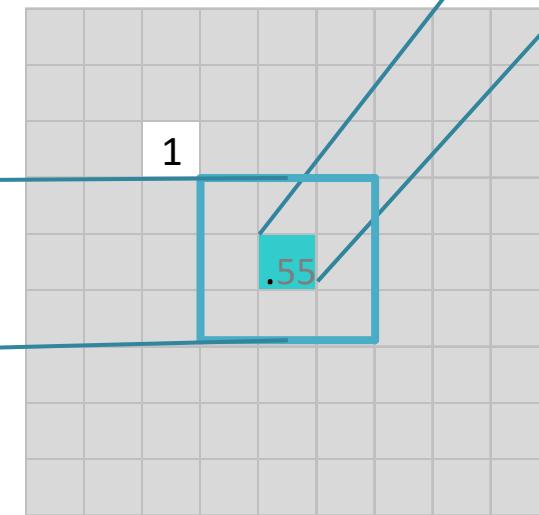
Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

1	1	-1
1	1	1
-1	1	1

$$\frac{1 + 1 - 1 + 1 + 1 + 1 - 1 + 1 + 1}{9} = .55$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



Convolution: Trying every possible match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

Convolution: Trying every possible match

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



$$\begin{matrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{matrix} =$$

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	1
-1	1	-1
1	-1	1

=

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



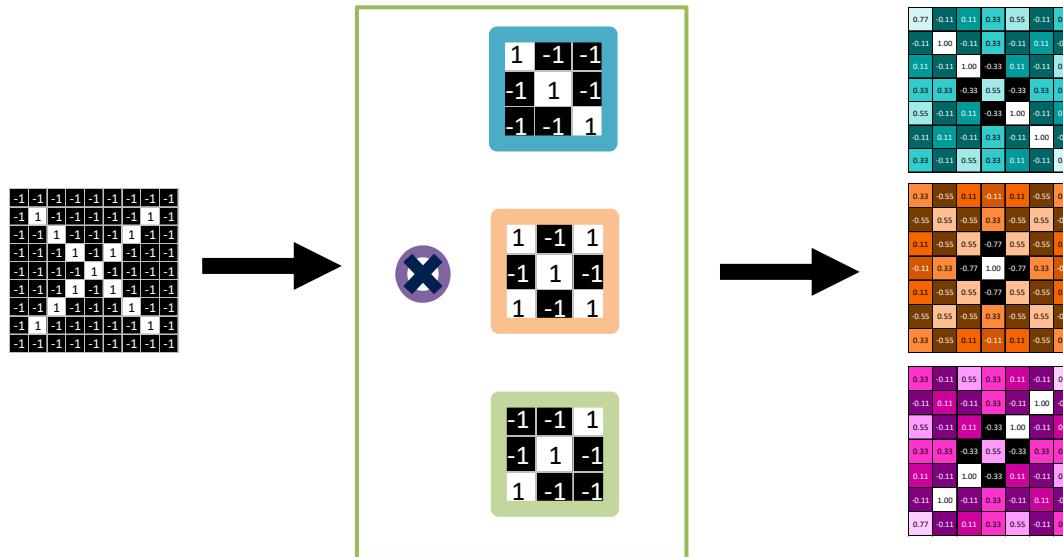
-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

Convolution layer

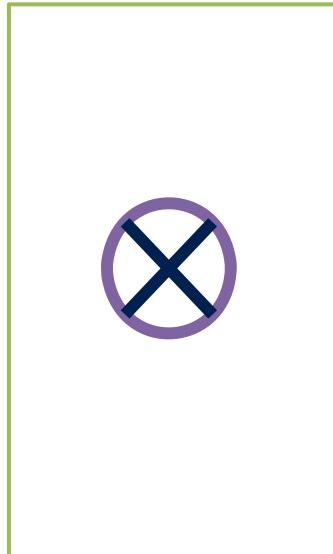
One image becomes a stack of filtered images



Convolution layer

One image becomes a stack of filtered images

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	0.33	1.00	0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	0.55	0.55	-0.55
0.11	-0.55	0.55	0.77	0.55	-0.55	0.11
0.11	0.33	-0.77	1.00	0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	0.55	0.55	-0.55
0.33	-0.55	0.11	0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	-0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	0.11	0.55
0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

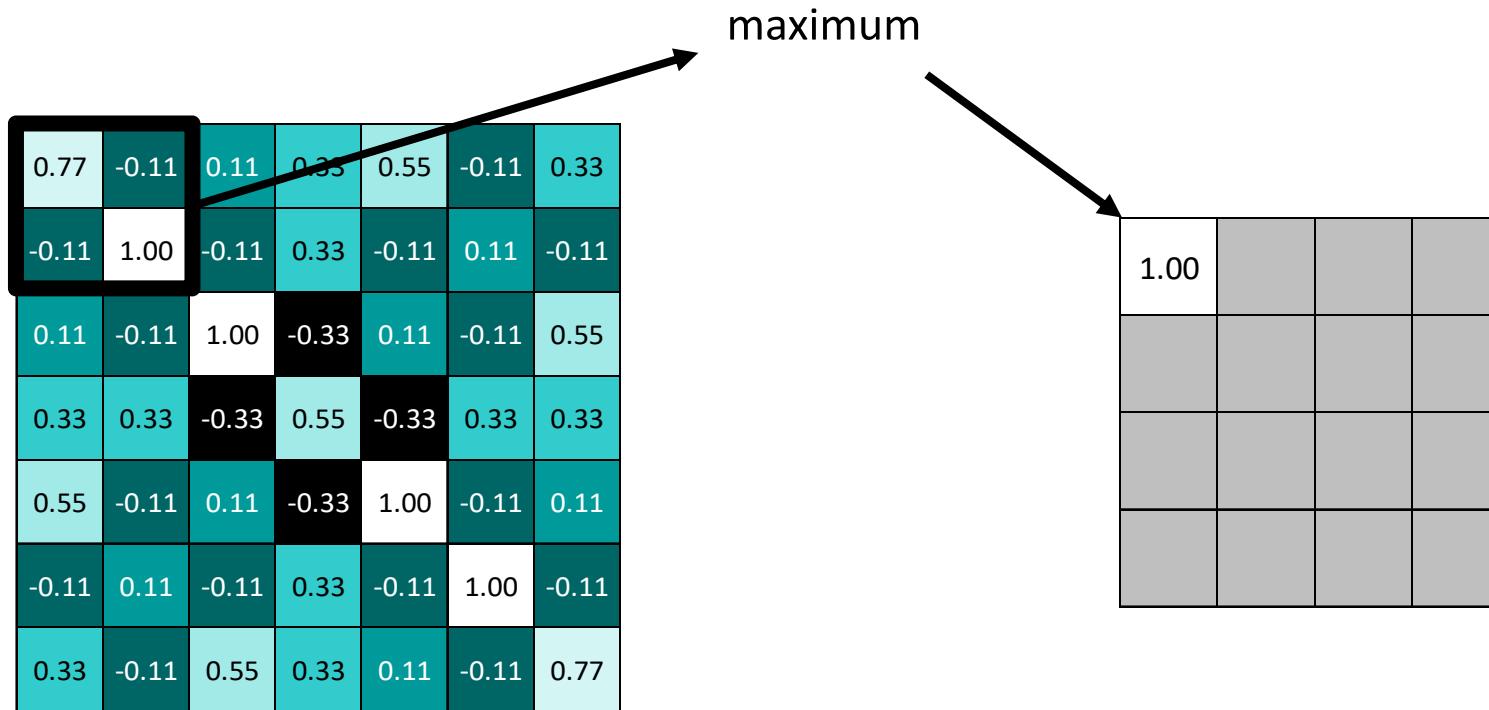
Pooling: Shrinking the image stack

Achieves a more abstract representation

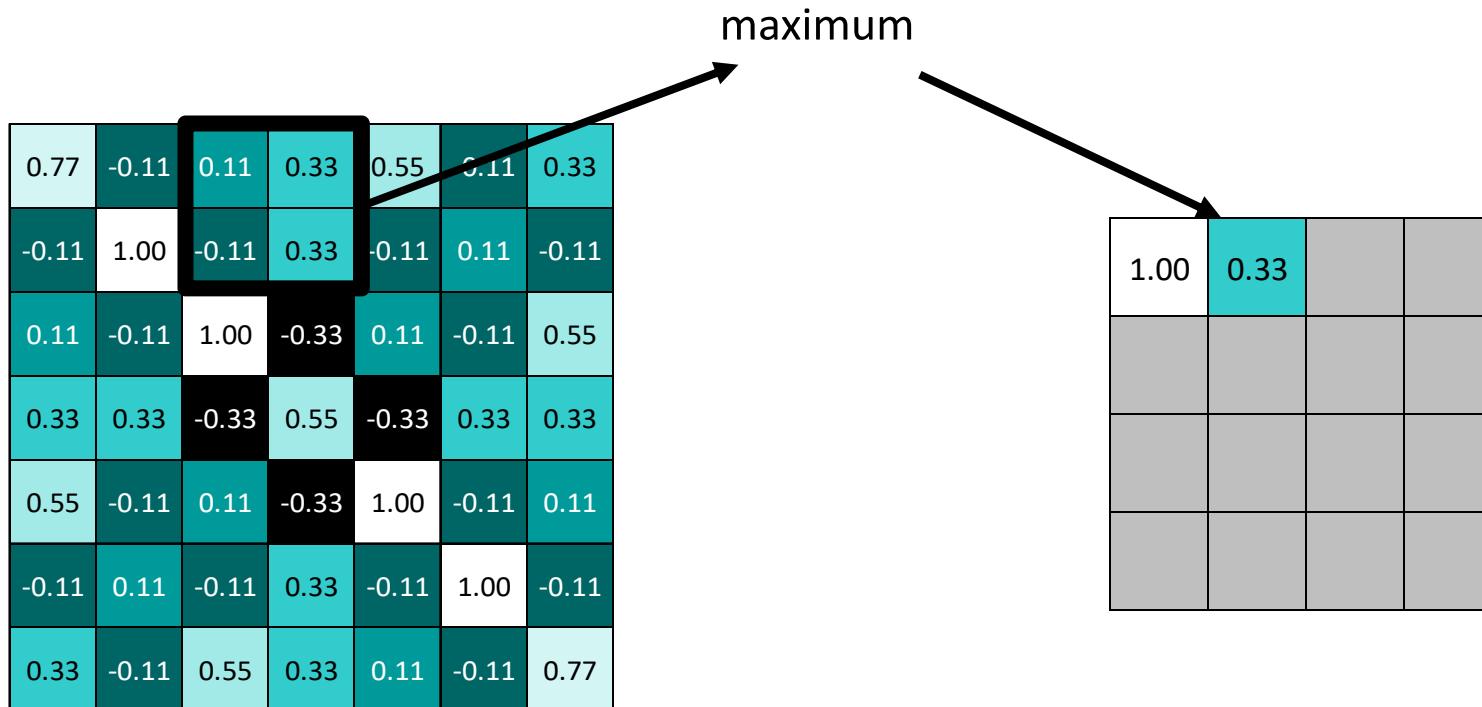
Improves invariance to geometric transformation

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.

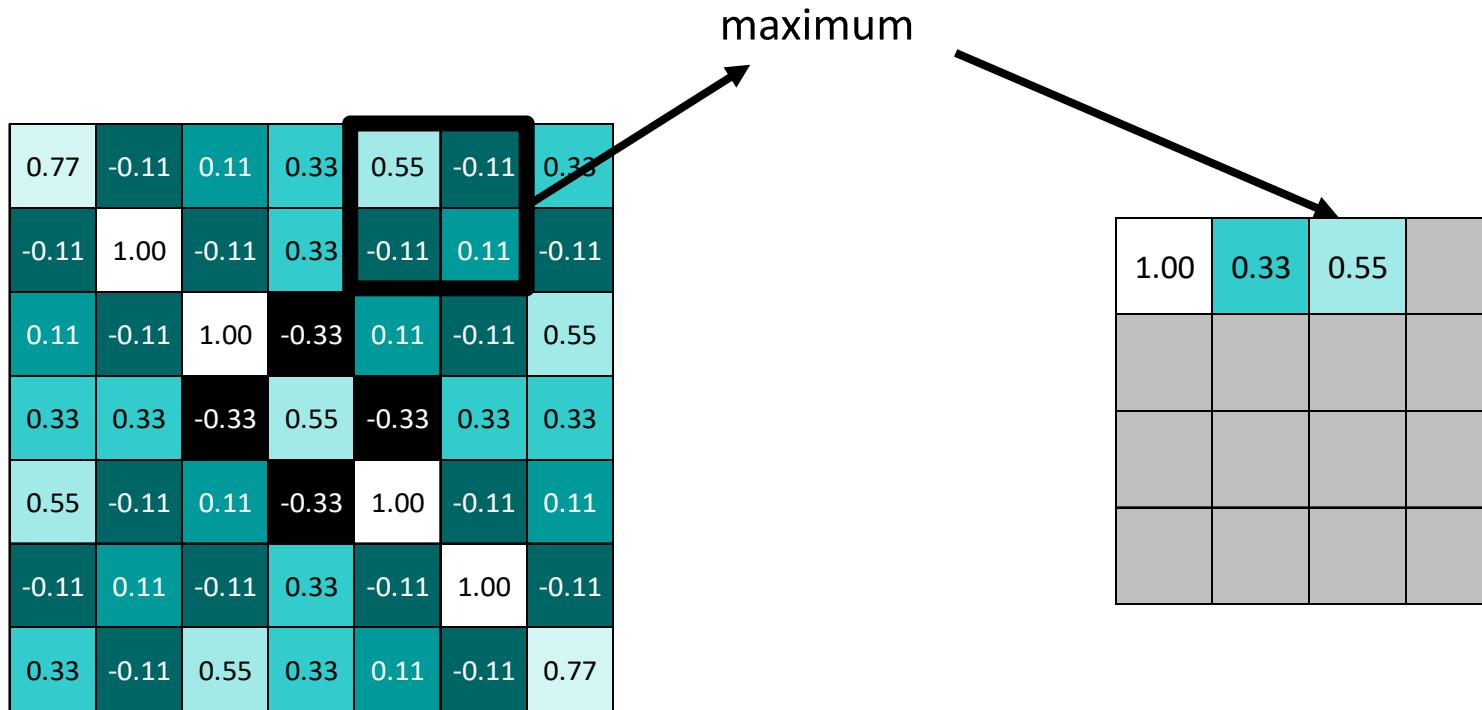
Pooling



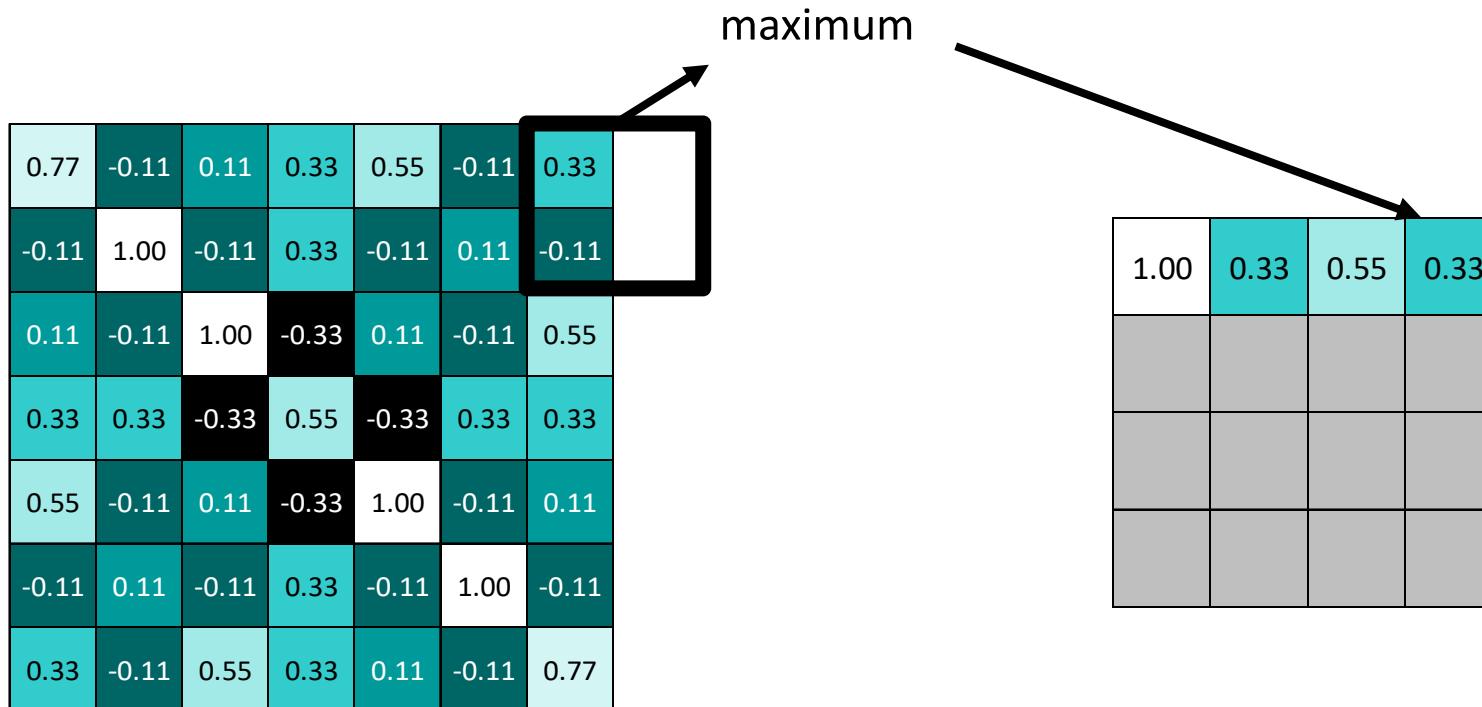
Pooling



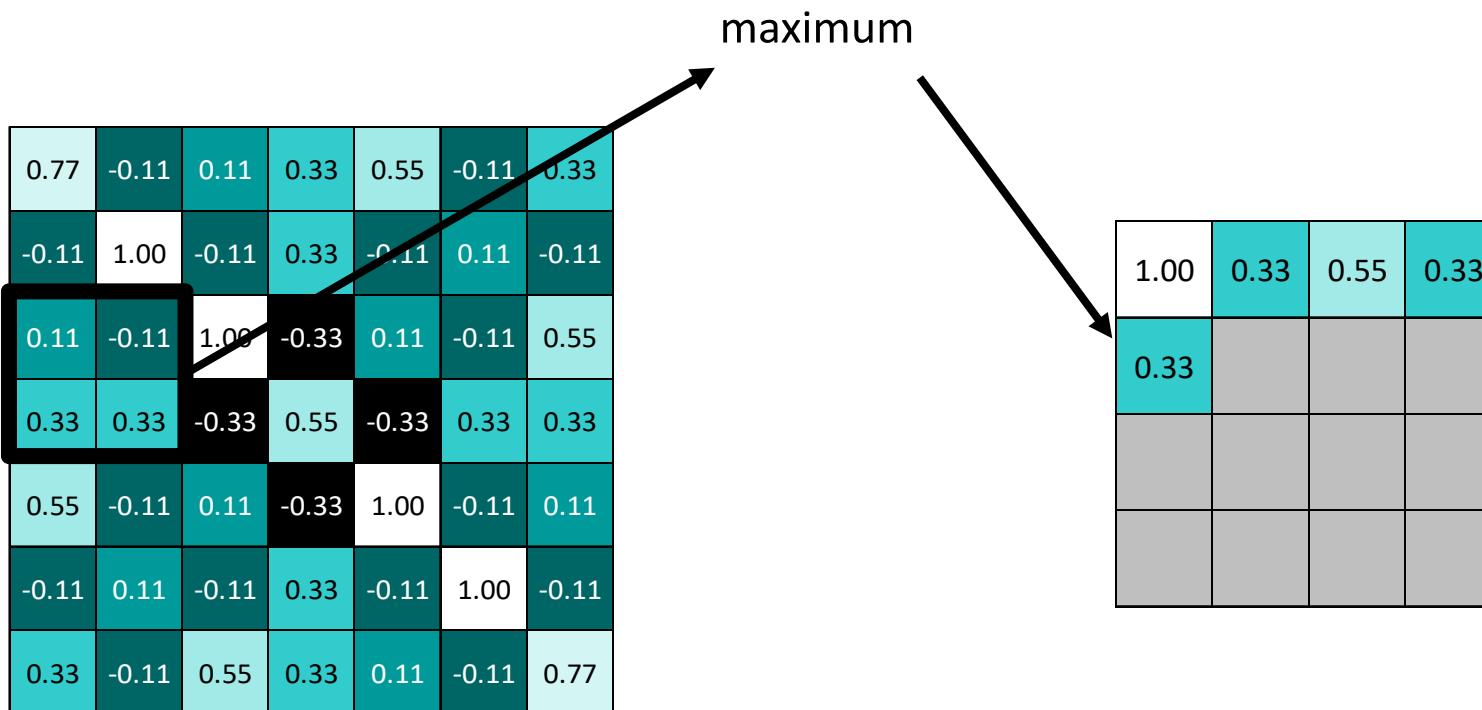
Pooling



Pooling



Pooling



Pooling

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

max pooling



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33



0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

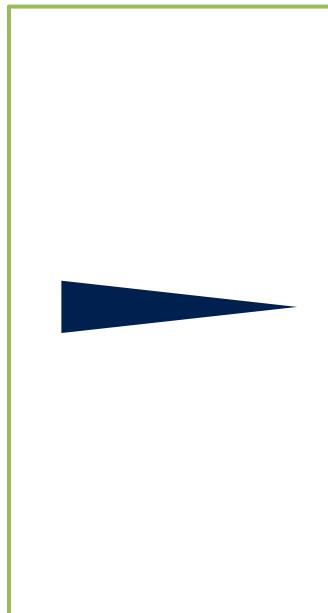
Pooling layer

A stack of images becomes a stack of smaller images.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

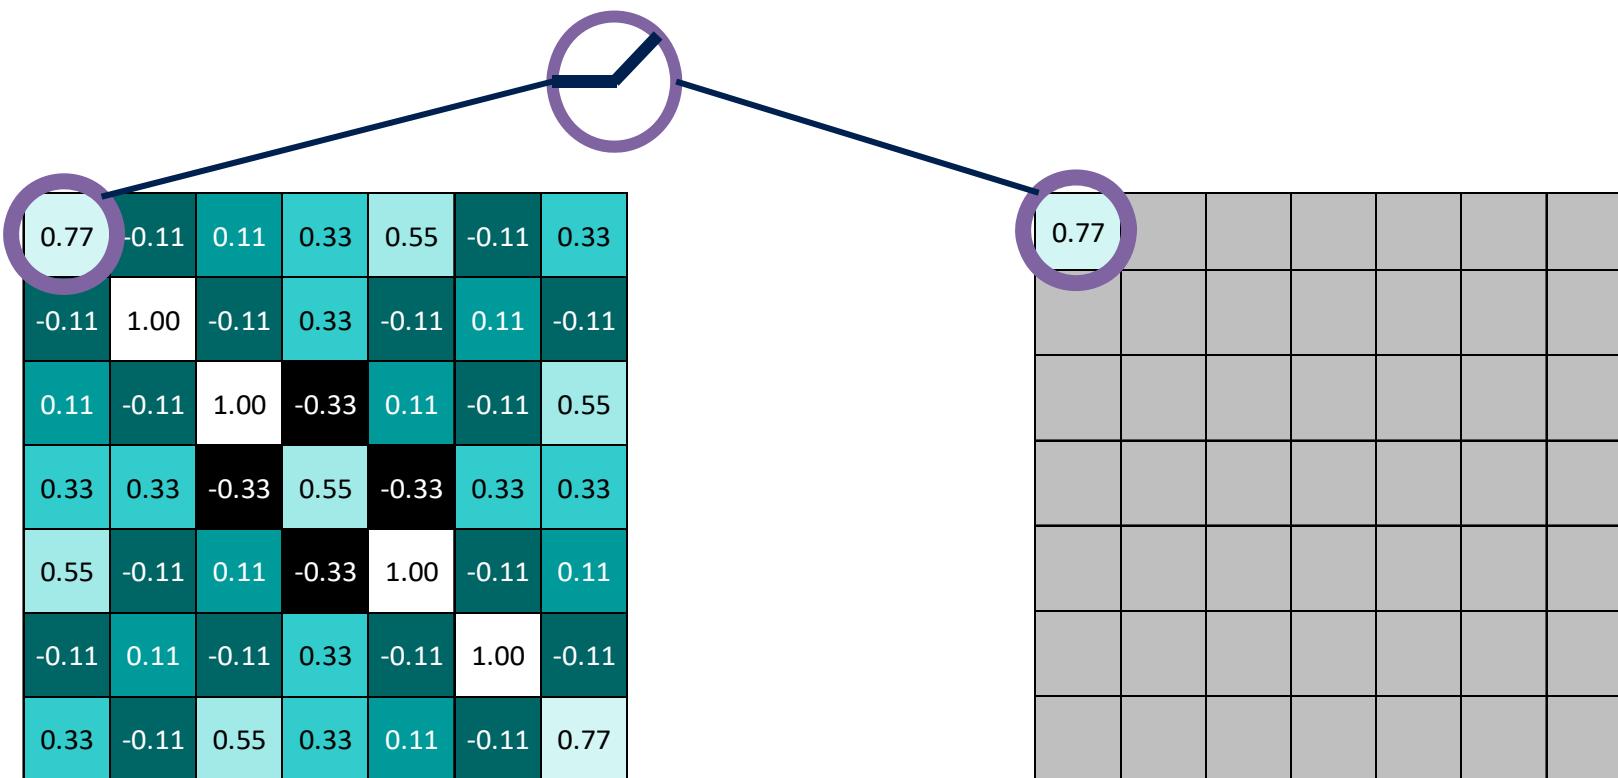
0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

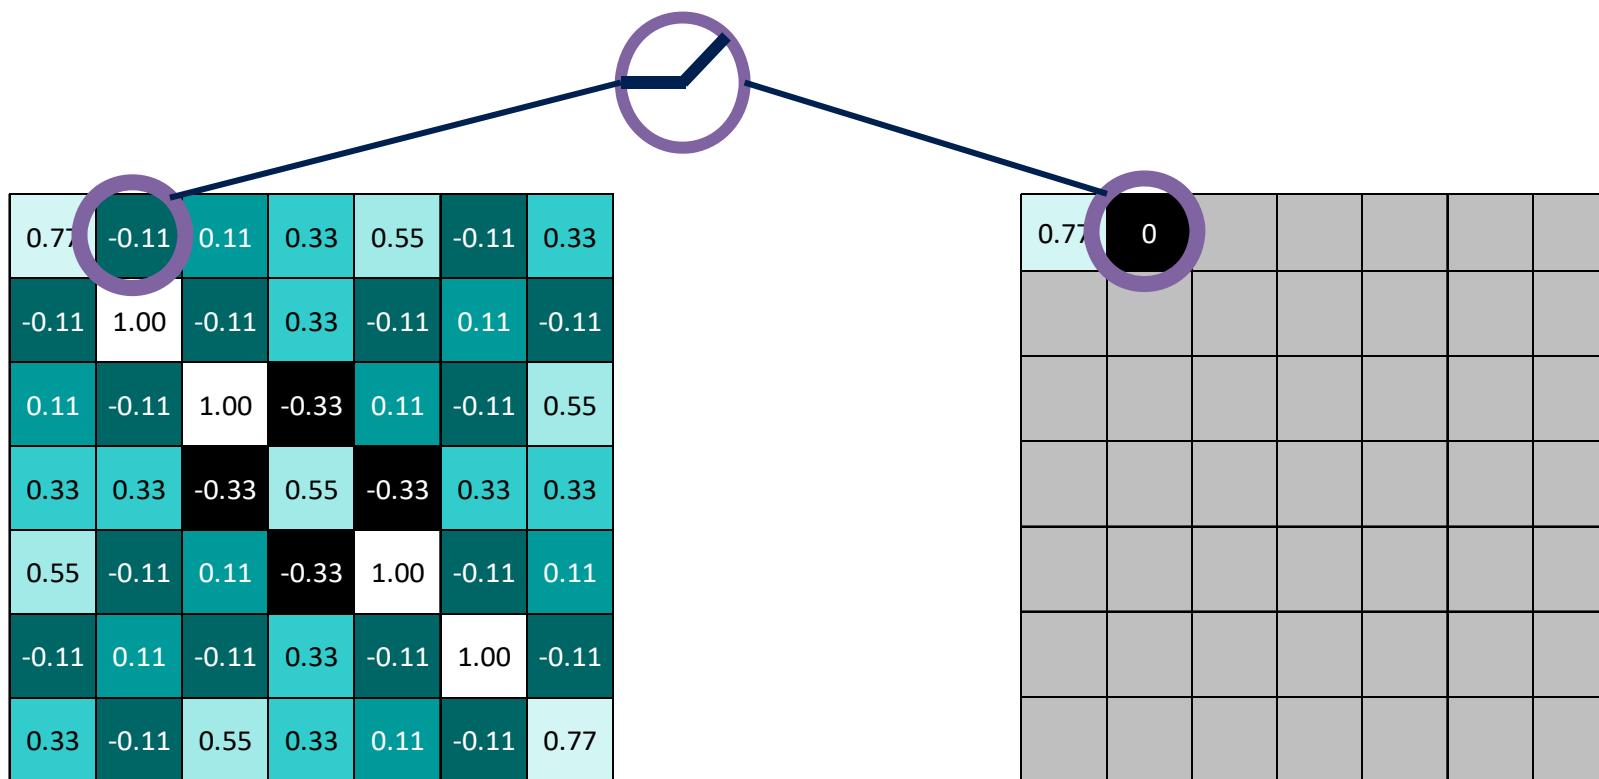
Normalization

- Keep the math from breaking by tweaking each of the values just a bit.
- Change everything negative to zero.

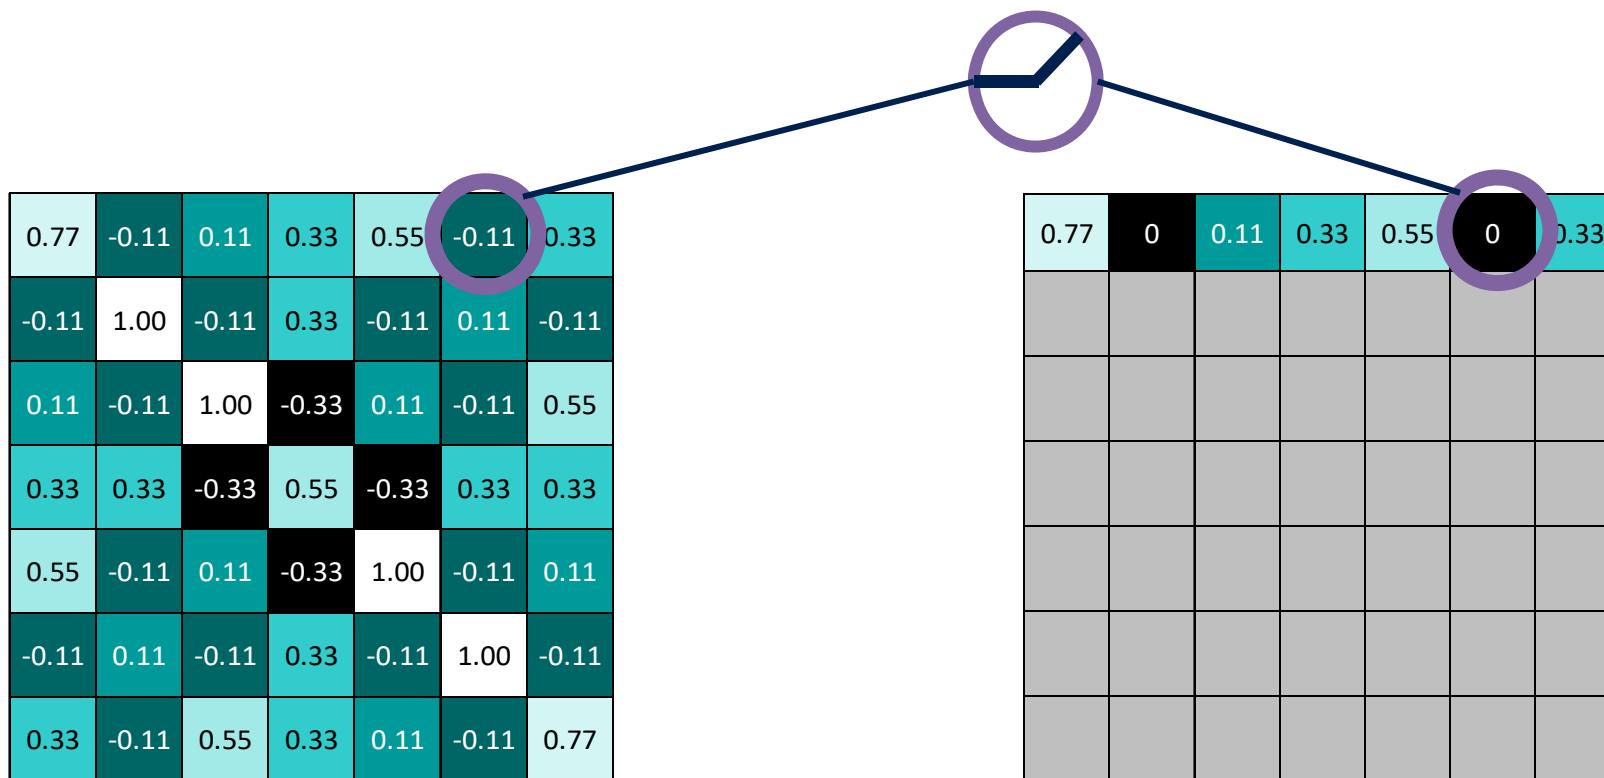
Rectified Linear Units (ReLUs)



Rectified Linear Units (ReLUs)



Rectified Linear Units (ReLUs)



Rectified Linear Units (ReLUs)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

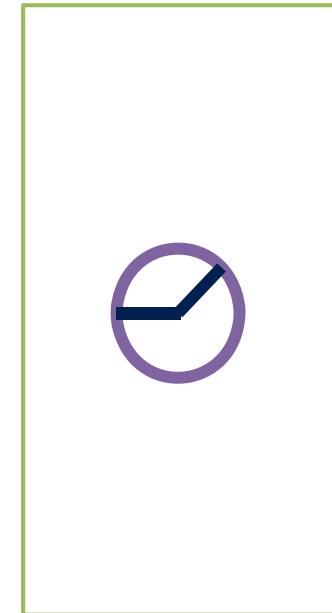
ReLU layer

A stack of images becomes a stack of images with no negative values.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	0.55	0.33	0.55	0.55	0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



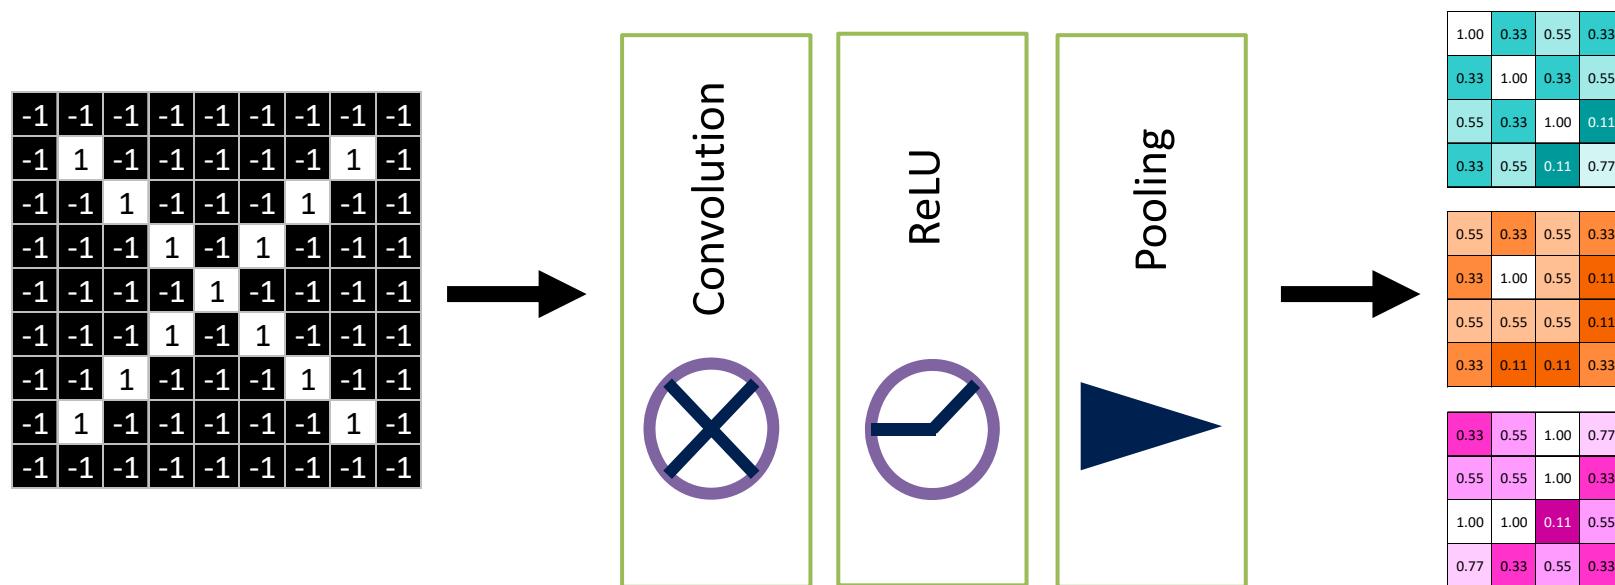
0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

0.33	0	0.11	0	0.11	0	0.33
0	0.55	0	0.33	0	0.55	0
0.11	0	0.55	0	0.55	0	0.11
0	0.33	0	1.00	0	0.33	0
0.11	0	0.55	0	0.55	0	0.11
0	0.55	0	0.33	0	0.55	0
0.33	0	0.11	0	0.11	0	0.33

0.33	0	0.55	0.33	0.11	0	0.77
0	0.11	0	0.33	0	1.00	0
0.55	0	0.11	0	1.00	0	0.11
0.33	0.33	0	0.55	0	0.33	0.33
0.11	0	1.00	0	0.11	0	0.55
0	1.00	0	0.33	0	0.11	0
0.77	0	0.11	0.33	0.55	0	0.33

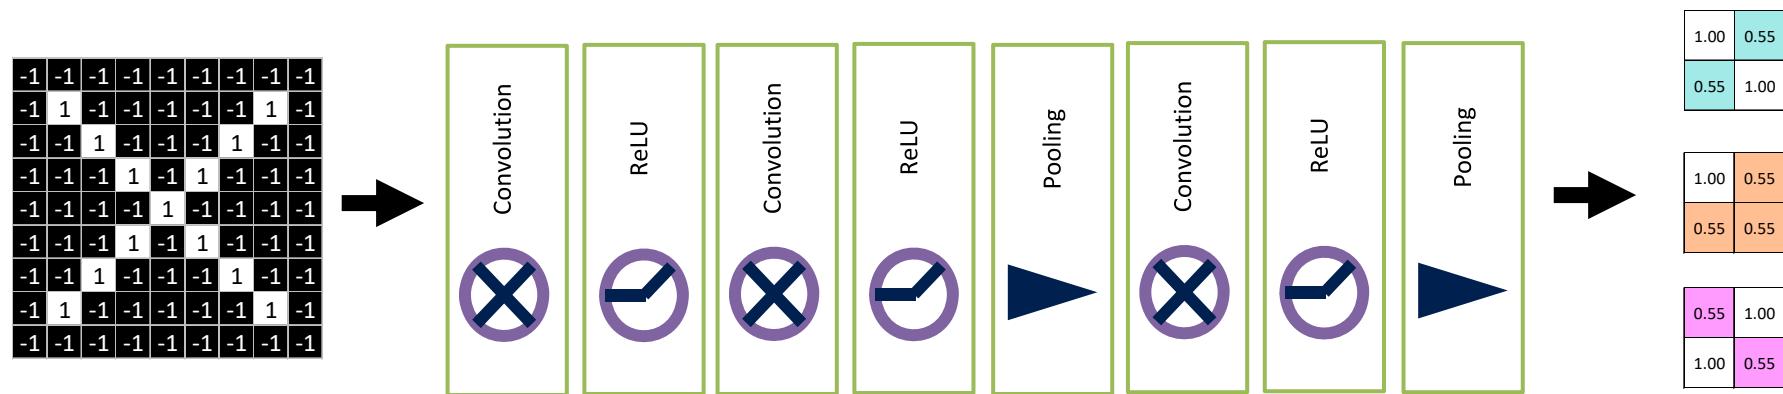
Layers get stacked

The output of one becomes the input of the next.



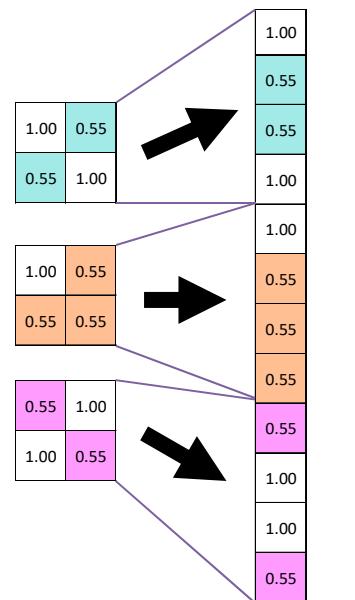
Deep stacking

Layers can be repeated several (or many) times.



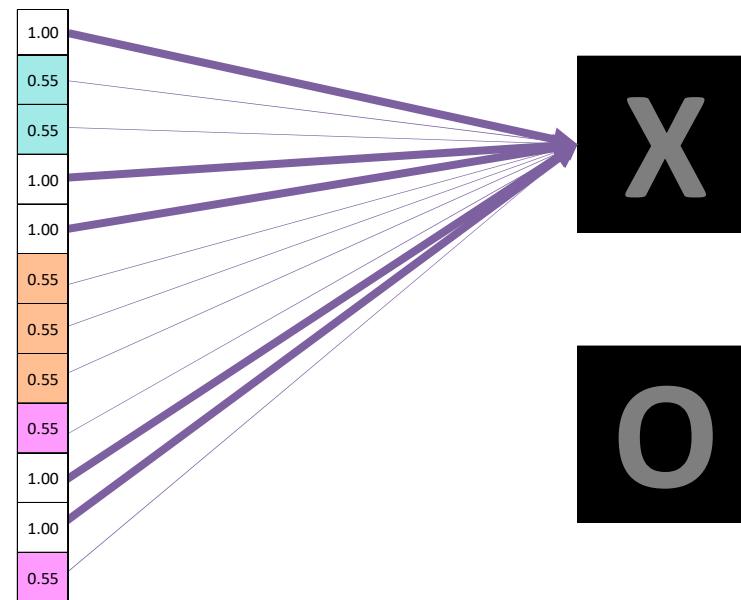
Fully connected layer

Every value gets a vote



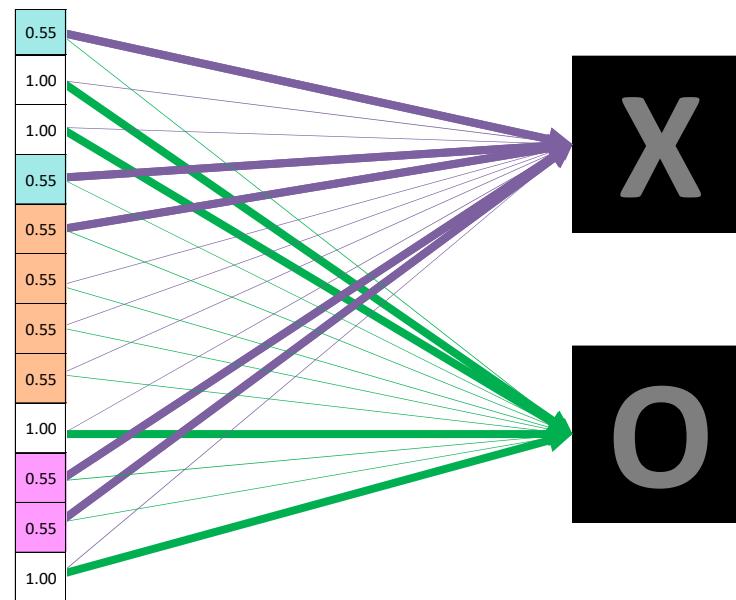
Fully connected layer

Vote depends on how strongly a value predicts X or O



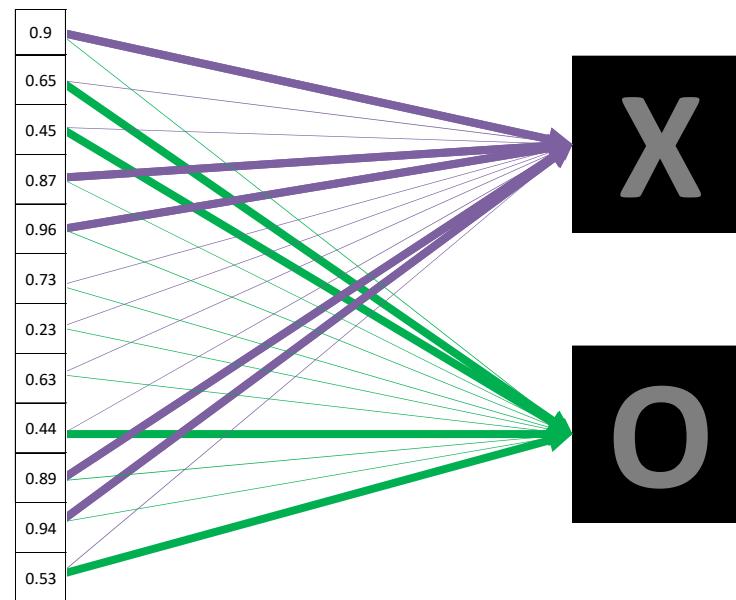
Fully connected layer

Vote depends on how strongly a value predicts X or O



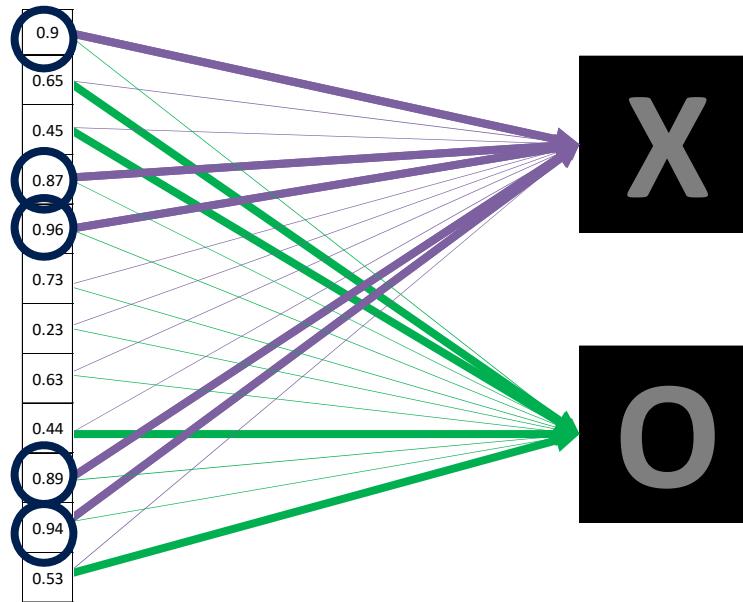
Fully connected layer

Future values vote on X or O



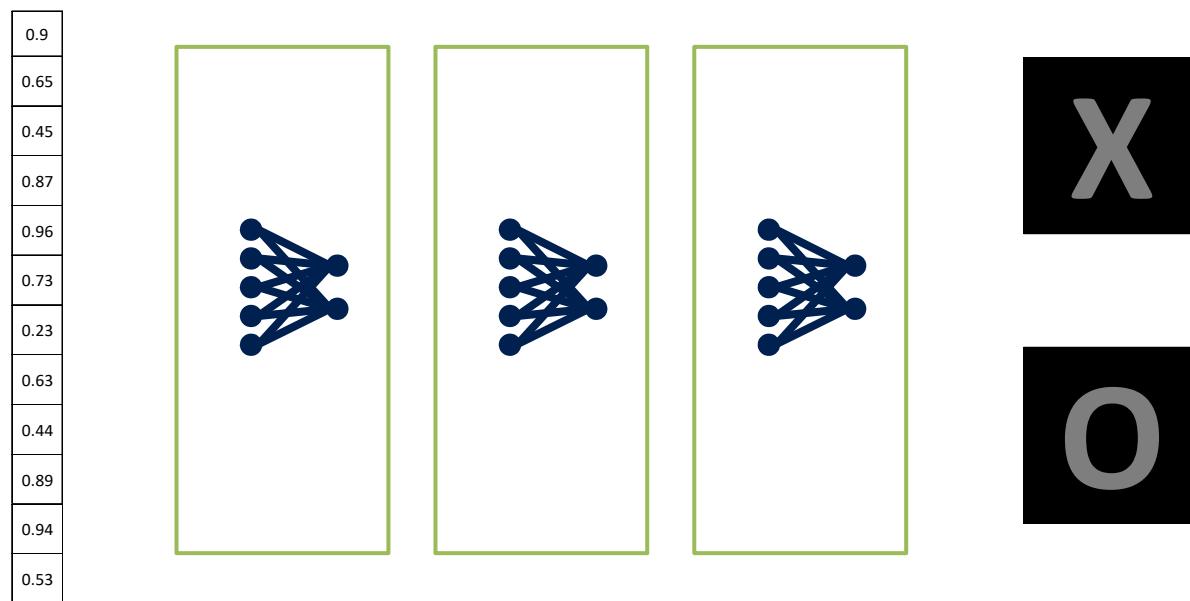
Fully connected layer

Future values vote on X or O



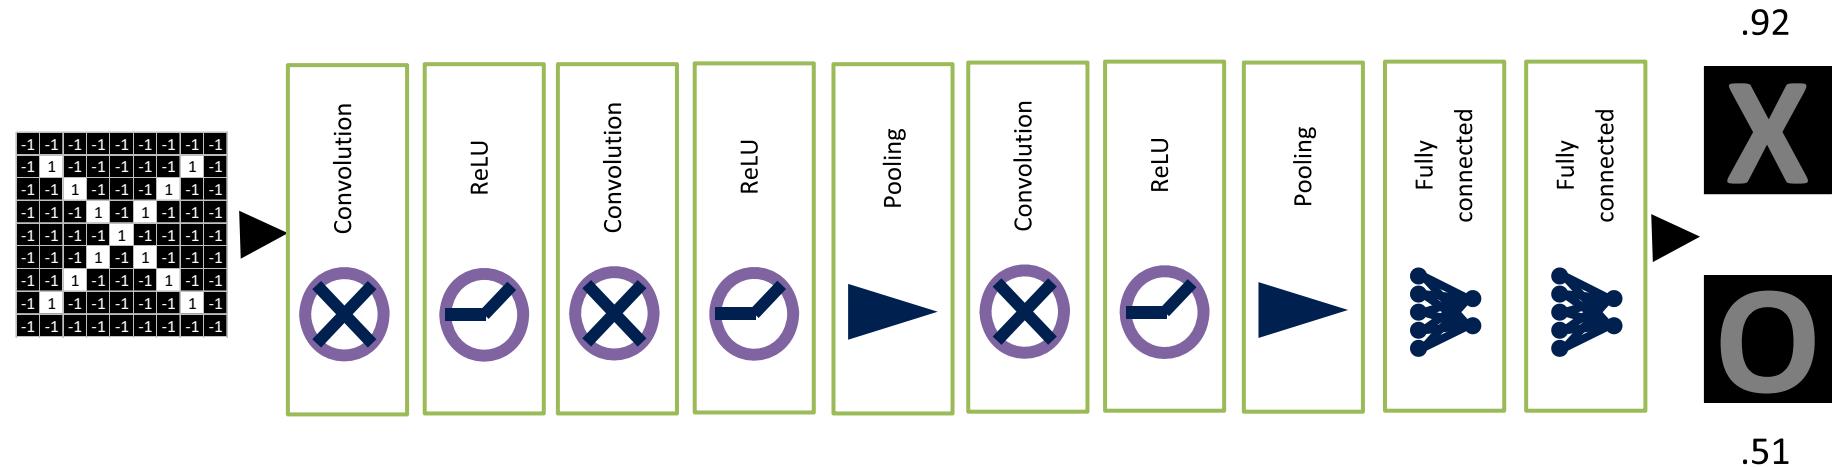
Fully connected layer

These can also be stacked.



Putting it all together

A set of pixels becomes a set of votes.



Learning

Q: Where do all the magic numbers come from?

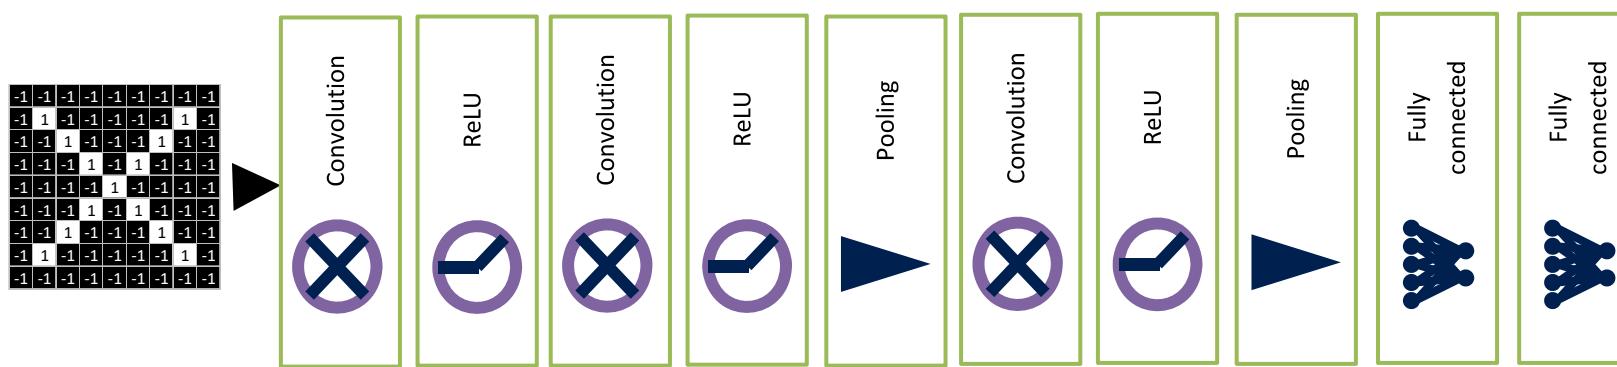
Features in convolutional layers

Voting weights in fully connected layers

A: Backpropagation

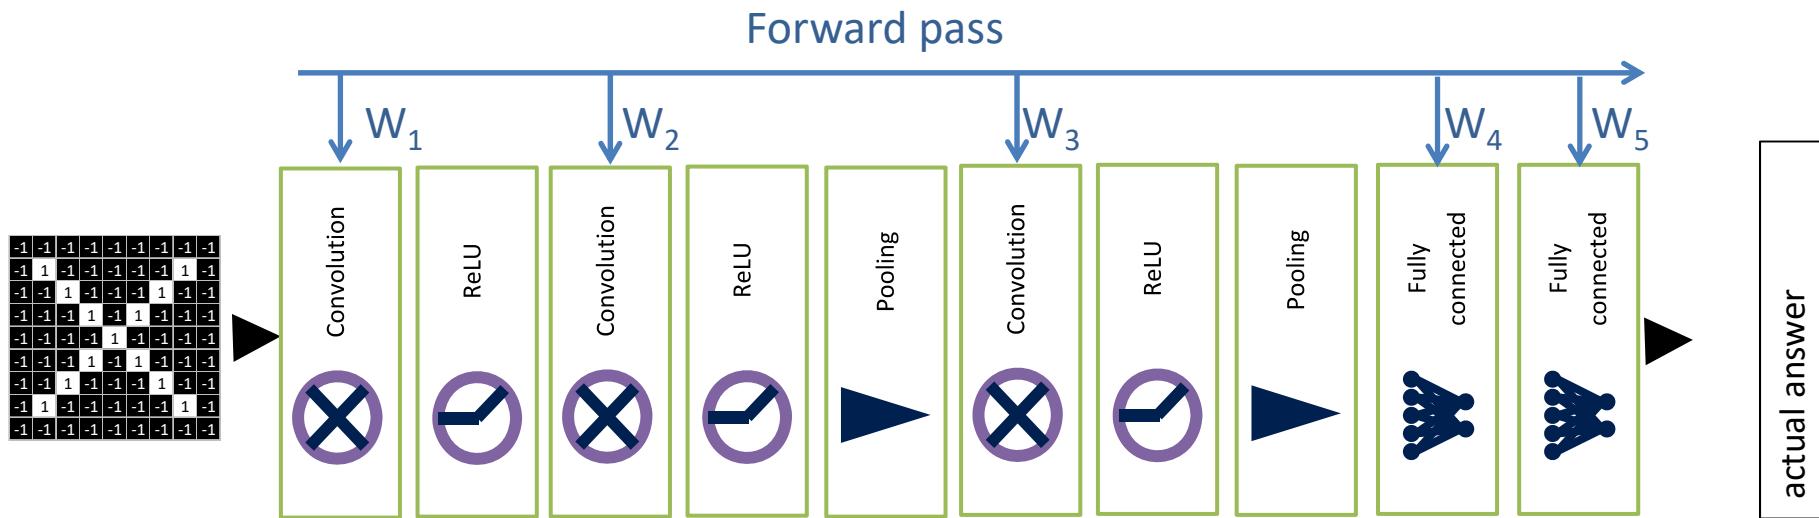
Backpropagation

In the first iteration W_i are set randomly



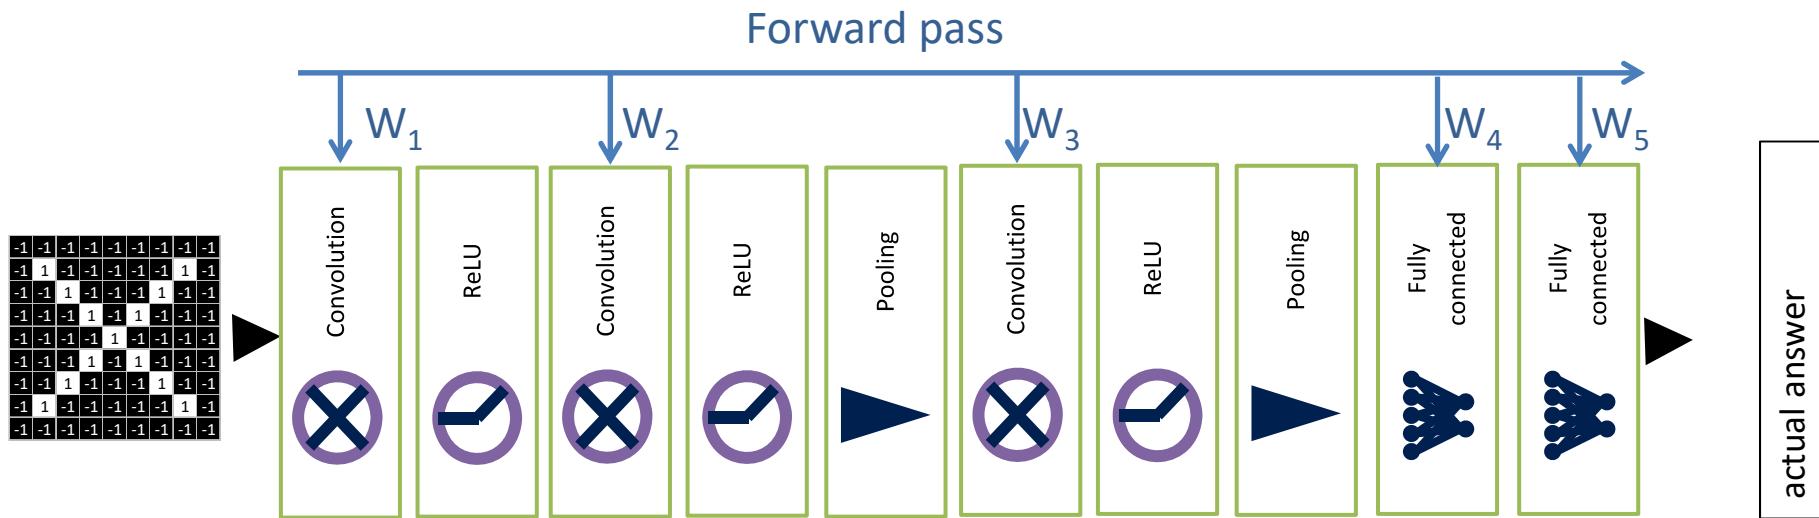
Backpropagation

In the first iteration W_i are set randomly



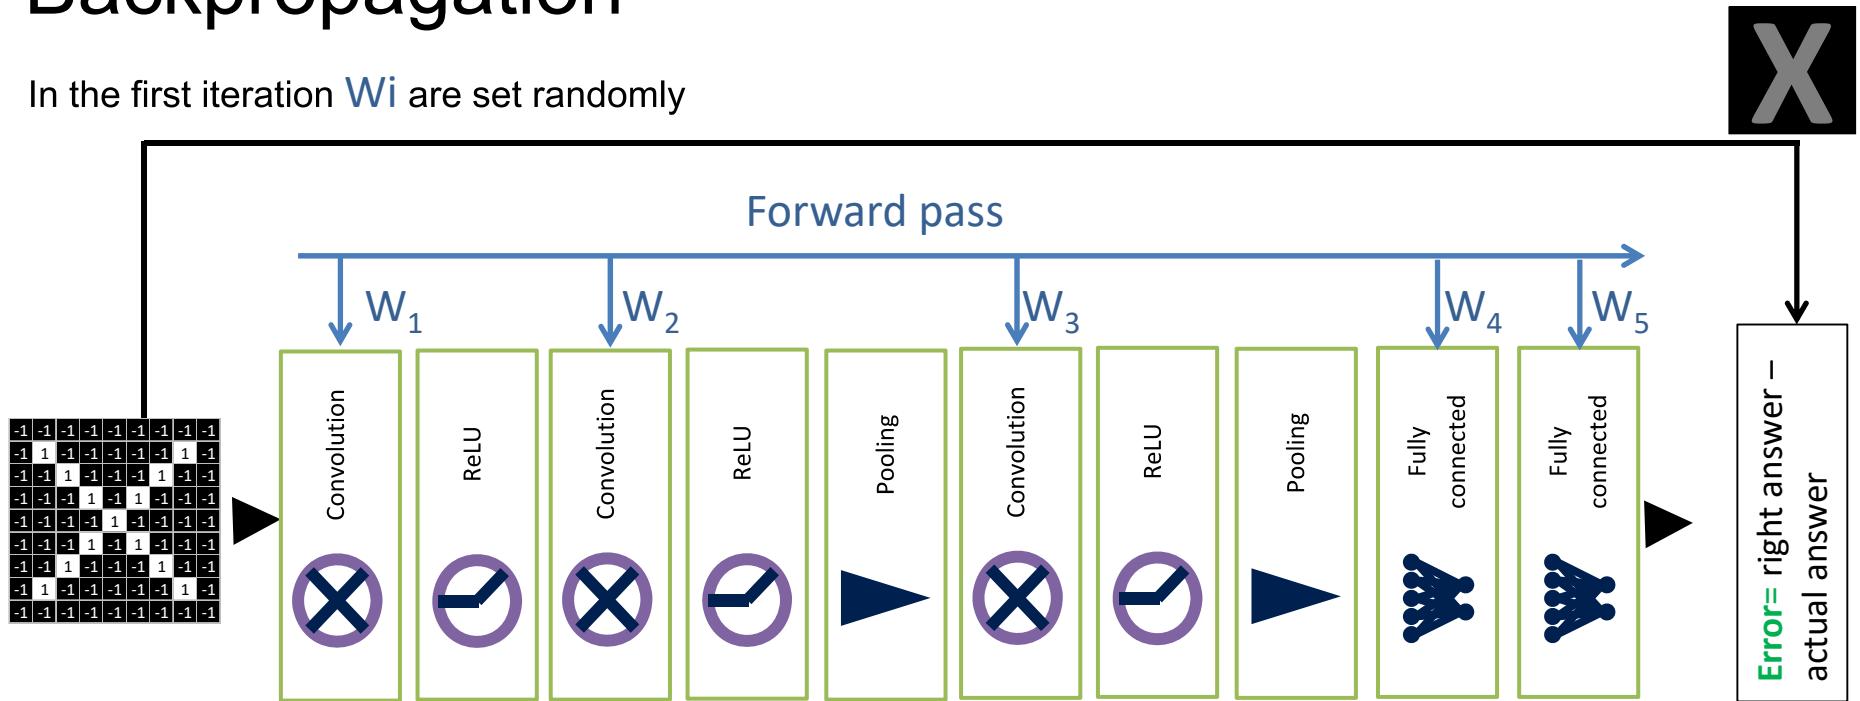
Backpropagation

In the first iteration W_i are set randomly



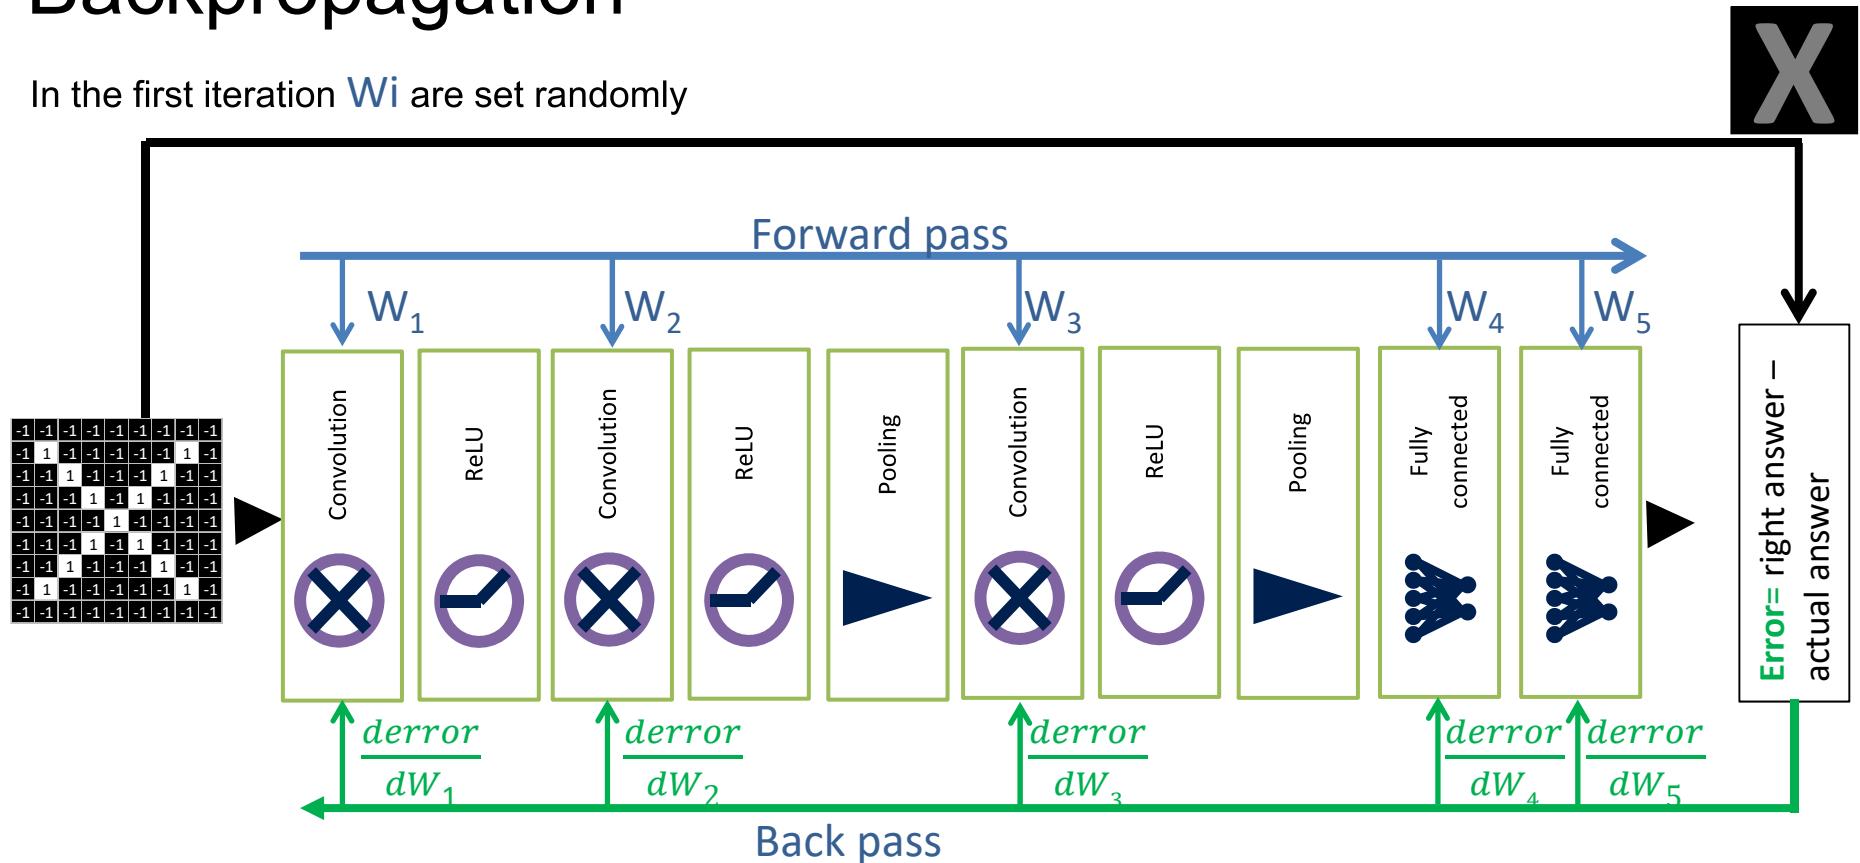
Backpropagation

In the first iteration W_i are set randomly



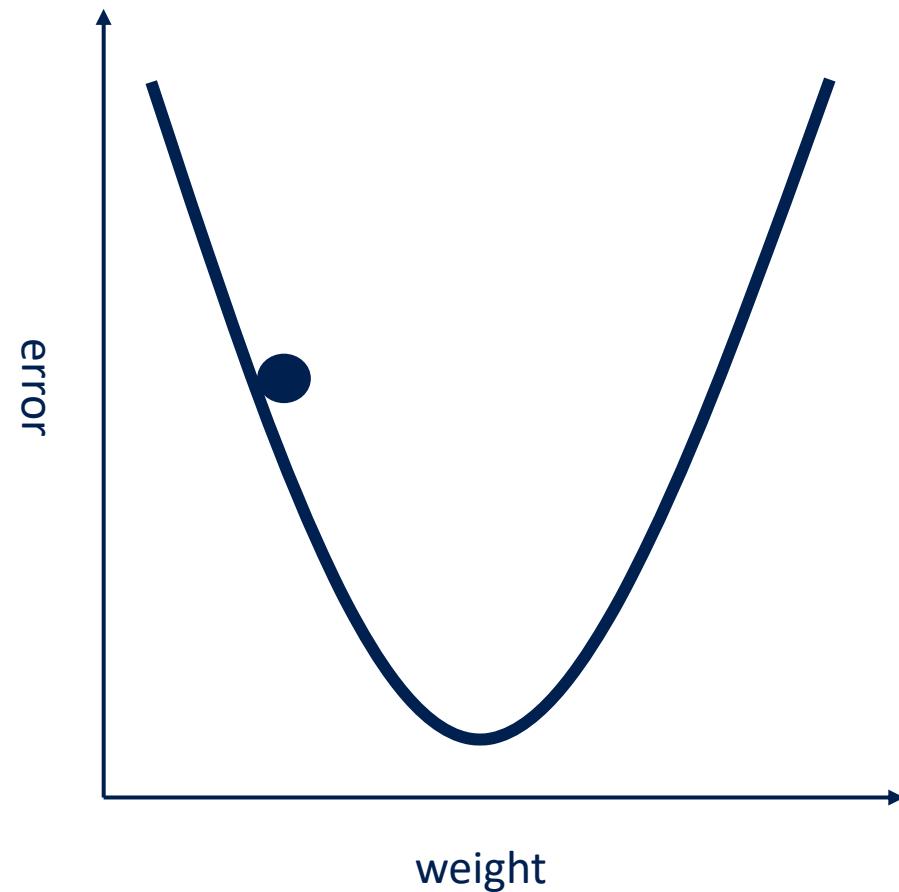
Backpropagation

In the first iteration W_i are set randomly



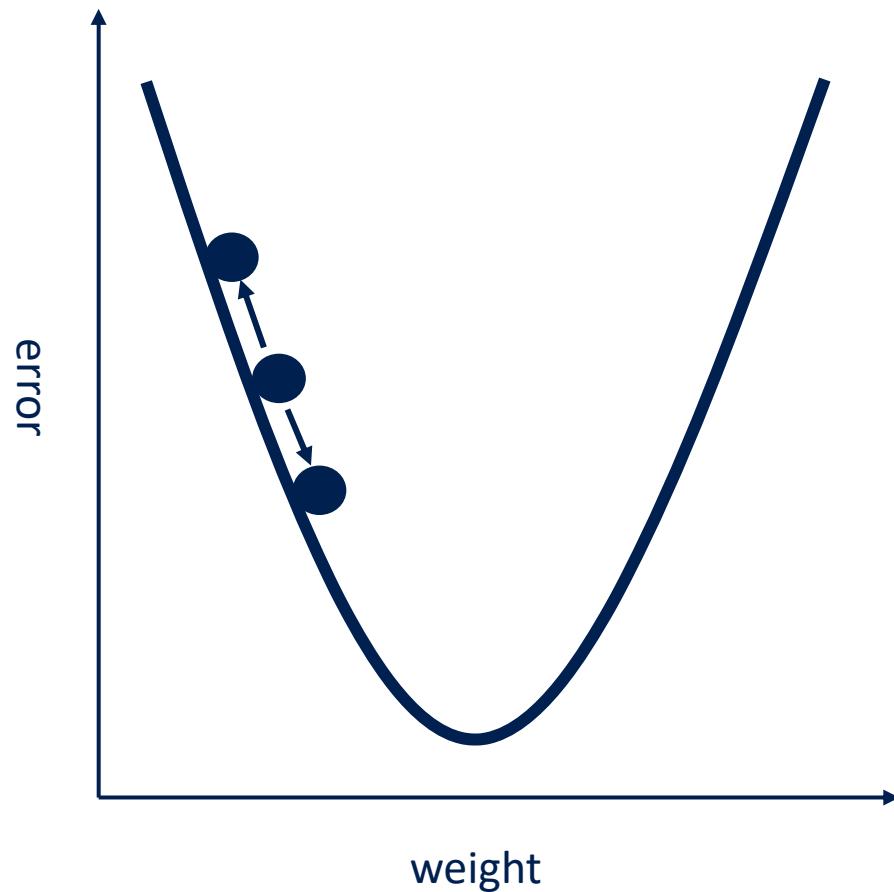
Gradient descent

For each feature pixel and voting weight,
adjust it up and down a bit and see how the
error changes.



Gradient descent

For each feature pixel and voting weight,
adjust it up and down a bit and see how the
error changes.



Limitations

- CNNs require large amount of data to get high accuracies

Practical solutions

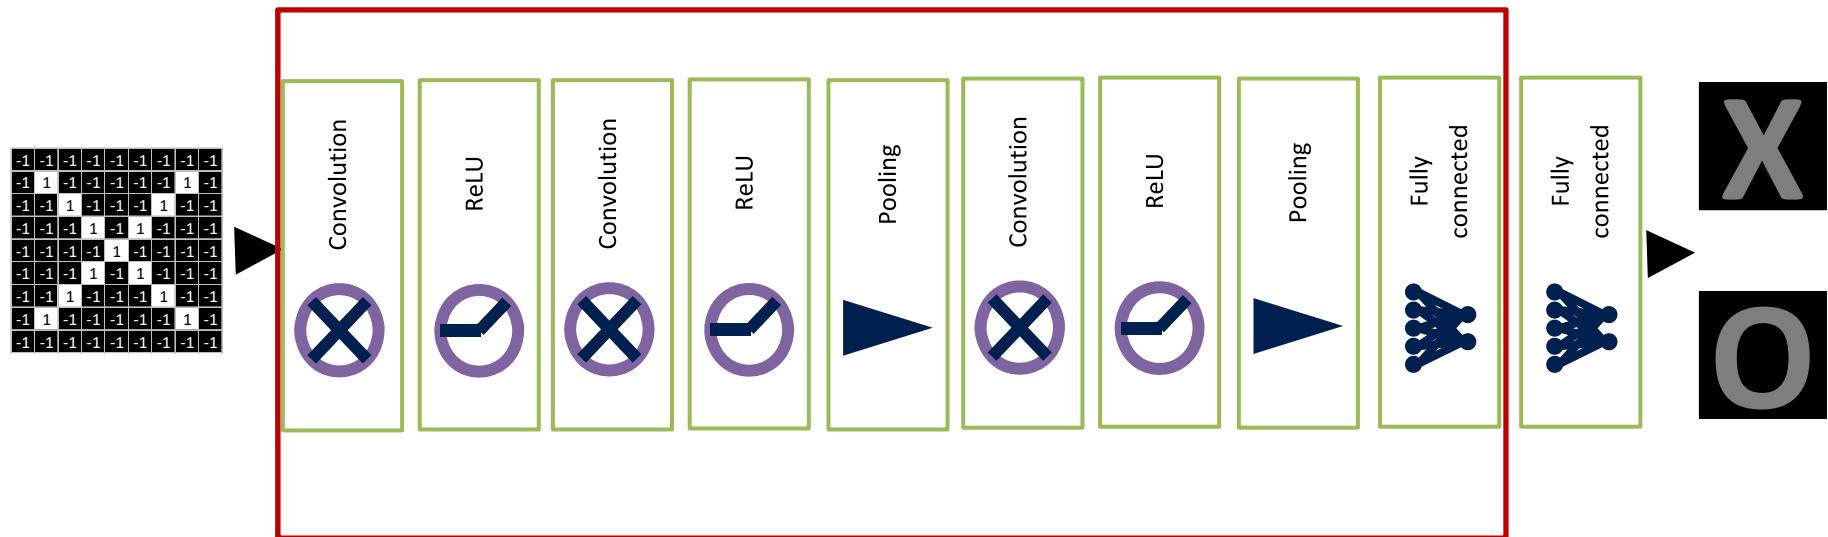
- Transfer learning
- Data augmentation

Outline

1. What are ANN?
2. What are CNNs?
3. How do CNNs work?
 - Convolution layers, pooling layers, FC layers, Backpropagation
4. Transfer learning
5. Data augmentation
6. One case study
7. CNNs for Detection
8. Two case studies

Transfer Learning

- ❑ Given a new problem & new data
- ❑ Instead of training from scratch, train only the last layers



Outline

1. What are ANN?
2. What are CNNs?
3. How CNNs work?
 - Convolution layers, pooling layers, FC layers, Backpropagation
4. Transfer learning
5. Data augmentation
6. Three case studies

Data augmentation

- ❑ Increase the training dataset volume artificially using transformations
- ❑ Objective: Improve model robustness



Original : Scaling Centring translation rotation etc

Outline

1. What are ANN?
2. What are CNNs?
3. How CNNs work?
 - Convolution layers, pooling layers, FC layers, Backpropagation
4. Transfer learning
5. Data augmentation
6. One case study
7. CNNs for Detection
8. Two case studies

Case study 1: Data-augmentation for CNNs using MNIST

- Objective: Analyze the benefit of data-augmentation and ensembles on CNNs
- Methodology:
 - MNIST (60.000 train + 30.000 test) & 10 classes
 - Three CNNs: LeNet, Network3, Dropconnect

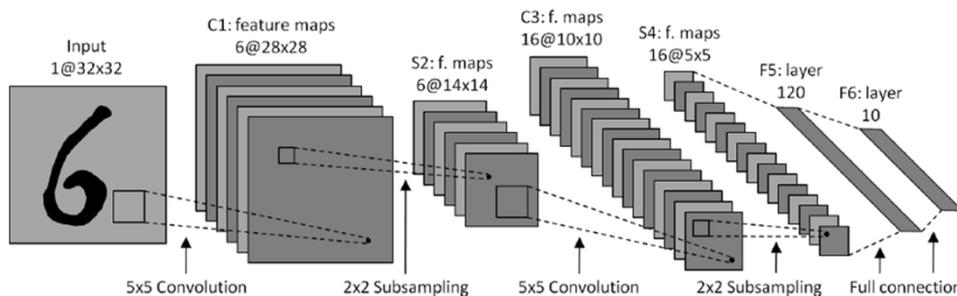


[A snapshot of image pre-processing for convolutional neural networks: case study of MNIST](#) S Tabik, D Peralta, A Herrera-Poyatos, F Herrera. International Journal of Computational Intelligence Systems 10, 555–568

Case study 1: Data-augmentation for CNNs using MNIST

Three Lenet like CNNs:

1. LeNet: 2 Conv + 2FC(ReLU)
2. Network3: 2 Conv + 3FC (ReLU+Dropout)
3. DropConnect: 2 Conv + 2FC (ReLU+Dropconnect)



Data-augmentation techniques

Dataset	Combination	# of training instances
1	Original	60,000
2	Centering	60,000
3	Elastic	300,000
4	Translation	300,000
5	Rotation	300,000
6	Elastic-centering	300,000
7	Rotation-centering	300,000
8	Translation-elastic	1,500,000
9	Translation-rotation	1,500,000
10	Rotation-elastic	1,500,000
11	Rotation-elastic-centering	1,500,000
12	Elastic-elastic	1,500,000

Case study 1: Data-augmentation for CNNs using MNIST

Test-set accuracies ($\frac{\text{correct predictions}}{\text{total test samples}}$)

Dataset	LeNet (10,000 iter.)				LeNet (50,000 iter.)			
	Average	Best	Epochs	Time(s)	Average	Best	Epochs	Time(s)
Original	99.08%	99.18%	10.67	267.91	99.05%	99.21%	213.33	1070.29
Centered	98.85%	99.06%	10.67	203.52	98.95 %	98.09%	213.33	926.38
Elastic	99.09%	99.19%	2.13	232.75	99.36%	99.44%	42.67	1065.38
Translation	99.09%	99.32%	2.13	268.75	99.30%	99.41%	42.67	1065.38
Rotation	99.05%	99.10%	2.13	268.03	99.25%	99.37%	42.67	1065.38
Elastic-centered	⁵ 99.17%	99.26%	2.13	267.20	99.27%	99.36%	42.67	925.51
Rotation-centered	98.90%	99.07%	2.13	232.73	99.19%	99.33%	42.67	950.38
Translation-elastic	⁴ 99.18%	99.32%	0.43	267.43	⁵ 99.39%	99.54%	8.53	1050.38
Translation-rotation	99.16%	99.40%	0.43	267.41	³ 99.40%	97.55%	8.53	1045.38
Rotation-elastic	¹ 99.31%	99.39%	0.43	268.14	¹ 99.47%	99.57%	8.53	1046.25
Rotation-elastic-centered	³ 99.19%	99.24%	0.43	232.30	² 99.43%	99.52%	8.53	925.68
Elastic-elastic	² 99.27%	99.45%	0.43	268.10	⁴ 99.40%	99.50%	8.53	1047.64

Case study 1: Data-augmentation for CNNs using MNIST

Test-set accuracies

Dataset	Network3(10 epochs)			Network3(20 epochs)		
	Average	Best	Time(s)	Average	Best	Time(s)
Original	99.01%	99.07%	124.45	99.25%	99.25%	205,21
Centered	98.73%	98.80%	118.32	98.97%	99.01%	196.92
Elastic	99.49%	99.54%	656,85	99.61%	99.67%	1200,33
Translation	99.49%	99.55%	631.53	99.59%	99.63%	1228,71
Rotation	99.44%	99.50%	636.25	99.44%	99.50%	1256,95
Elastic-centered	99.32%	99.39%	566.44	99.57%	99.60%	1109,43
Rotation-centered	98.88%	98.94%	569.04	99.31%	99.32%	1167,32
Translation-elastic	99.54%	99.57%	3647.78	99.58%	99.63%	7111,65
Translation-rotation	99.57%	99.61%	3650.66	99.58%	99.60%	7149,25
Rotation-elastic	99.62%	99.67 %	3642,85	99.67%	99.69%	6996,23
Rotation-elastic-centered	99.43%	99.51%	3054,43	99.51%	99.52%	6908,70
Elastic-elastic	99.65%	99.66%	3607.32	99.67 %	99.70 %	7189,16

Case study 1: Data-augmentation for CNNs using MNIST

Test-set accuracies

Dataset	DropConnet(100 epochs)			DropConnet(200 epochs)		
	Average	Best	Time(s)	Average	Best	Time(s)
Original	98.32%	98.83%	7803.43	98.98%	98.99%	18748.53
Centered	95.35%	94.46%	6659.31	95.13%	98.85%	18635.54
Elastic	99.33 %	99.35%	7512.25	99.36%	99.36%	18606.15
Translation	⁵ 99.43%	99.46%	7736.41	⁵ 99.47%	99.47%	18710.45
Rotation	99.18%	99.29%	7151.73	99.37%	99.47%	18729.29
Elastic-centered	96.58%	96.69%	6969.89	97.08%	97.09%	18661.80
Rotation-centered	98.30%	98.41%	6974.23	98.55%	98.63%	18668.05
Translation-elastic	99.40%	99.57%	7162.37	³ 99.58%	99.58%	18745.93
Translation-rotation	² 99.57%	99.59%	7410.32	¹ 99.69%	99,69%	18772.40
Rotation-elastic	³ 99.54%	99,60%	7397.40	⁴ 99.56%	99.56%	18724.38
Rotation-elastic-centered	⁴ 99.47%	99.49%	7803.73	99.44%	99.46%	18220.50
Elastic-elastic	199,58%	99.59%	7911.30	² 99.59%	99.61%	18712.22

Case study 1: Data-augmentation for CNNs using MNIST

Results Esembles using the most voted strategy

	LeNet(500 neurons)		Network3		DropConnect	
	10,000 iter	50,000 iter	10 epochs	20 epochs	100 epochs	200 epochs
Ensemble-5	99,55%	99,57%	99,72%	99,69%	99,72%	99,66%
Ensemble-3	99,43%	99,54%	99,69%	99,67%	99,69%	99,68%

Error : 0.28% versus state of the art ensemble 0.23%

Case study 1: Data-augmentation for CNNs using MNIST

Results:The 28 misclassified characters

(2) 1 (5) 3 * (1) 7 (8) 9 (6) 5 (9) 4 (7) 1



(6) 0 (2) 7 (5) 3 * (1) 7 * (9) 5 (9) 7 (8) 3



(5) 3 (9) 4 * (9) 8 (4) 9 * (5) 3 * (5) 3 * (6) 1 *



(9) 4 * (1) 2 (5) 3 * (5) 6 (4) 9 * (2) 0 (5) 3 *



(7) 9 (6) 0 * (2) 7 * (1) 7 (2) 7 (6) 1 (3) 5



(6) 1 * (4) 9 (6) 0 * (5) 0 (6) 8 (7) 8 (2) 7 *



(9) 4 * (5) 3 * (3) 8 (7) 1 (8) 5 * (4) 9 (5) 6 *



(9) 5 (9) 8 * (7) 2 (5) 3 * (1) 7 (8) 5 * (5) 6 *



ensemble-5 (Network3)

ensemble-5 (DropConnect)

Case study 1: Data-augmentation for CNNs using MNIST

The 13 handwritten digits
misclassified by ensemble-5 of
DropConnet
and Network3

5, Net(3), DC(3) 1, Net(7), DC(7) 9, Net(4), DC(4) 5, Net(3), DC(3)



4, Net(9), DC(9) 5, Net(3), DC(3) 6, Net(1), DC(1) 6, Net(0), DC(0)



2, Net(7), DC(7) 9, Net(4), DC(8) 5, Net(3), DC(3) 8, Net(5), DC(5)



5, Net(6), DC(6)



Detection with CNNs

Siham Tabik

siham@ugr.es

Outline

- What is detection?
- How detection models work?
- The state-of-the art detection models
- Case study 1
- Case study 2

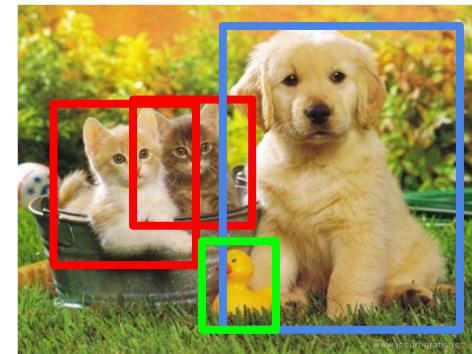
By the way, what is object detection?

Three classes: cat, dog and duck.

Given a test image:



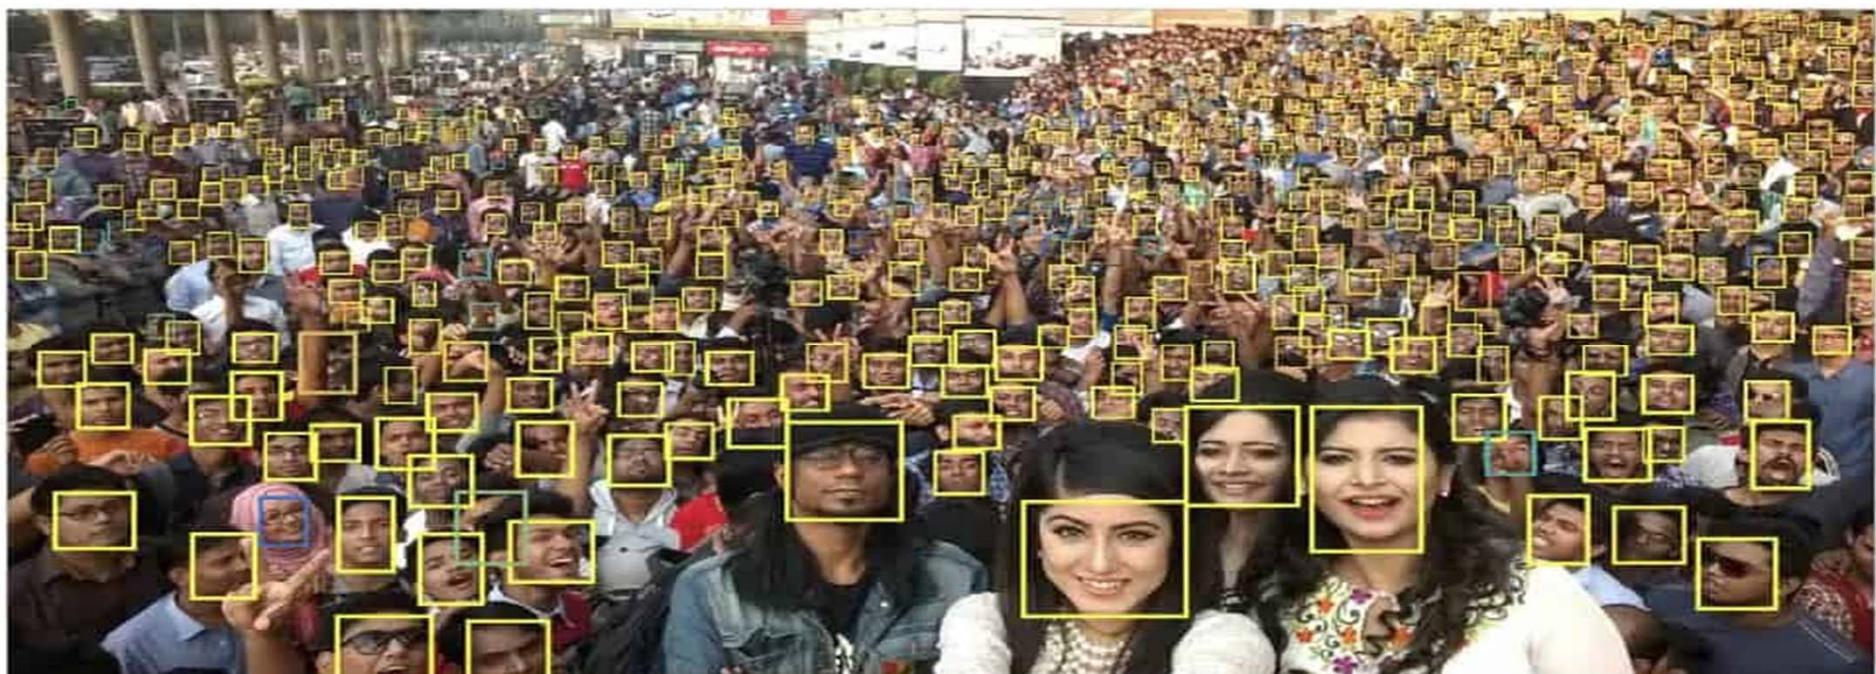
Object Detection



CAT, DOG, DUCK

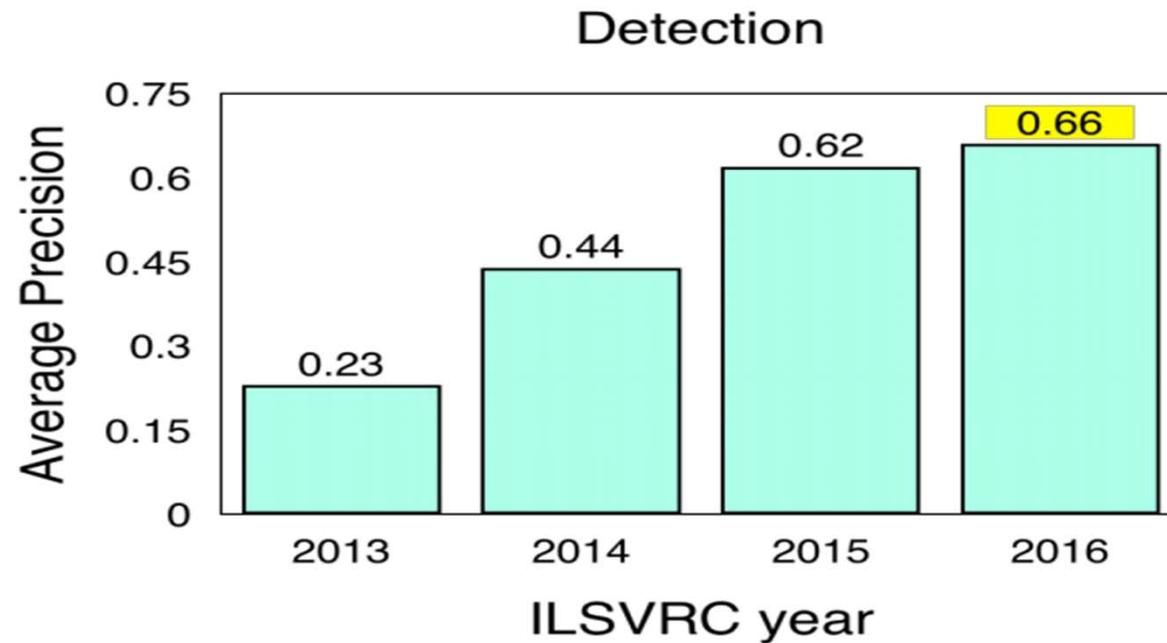
Object detection

An example of face detection: Object Detection of a single class.



Results of the ILSVRC object detection task over four years

Database: 200 object classes 527,982 images



How object detection works?

- Object detection methods reformulate the detection problem into a classification problem
- Combine a selective search method with a classification model

Detection as Classification



CAT? NO

DOG? NO

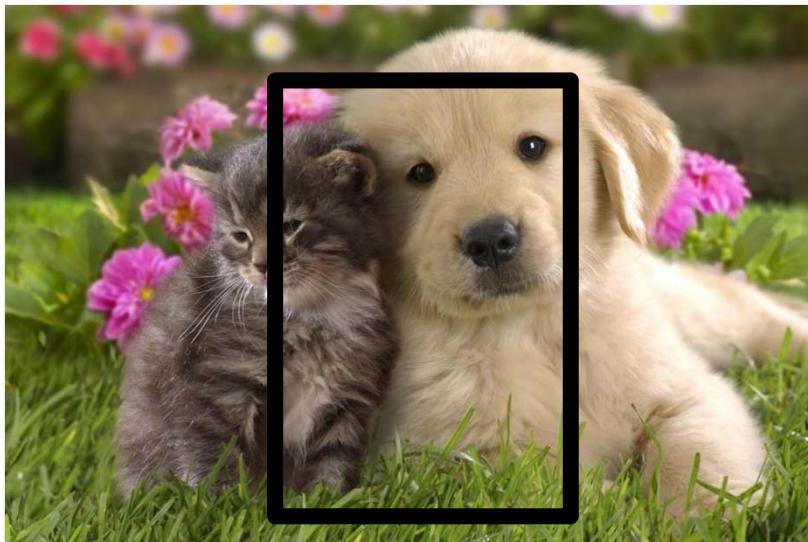
Detection as Classification



CAT? YES!

DOG? NO

Detection as Classification



CAT? NO

DOG? NO

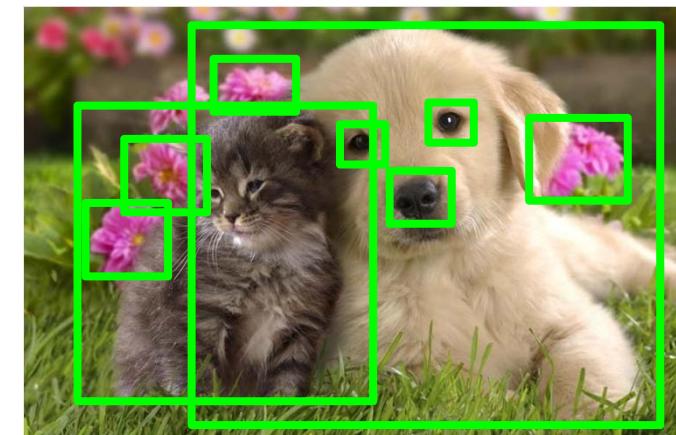
Detection as Classification

Problem: Need to test many positions and scales, and use a computationally demanding classifier (CNN)

Solution: Only look at a tiny subset of possible positions

Region Proposals

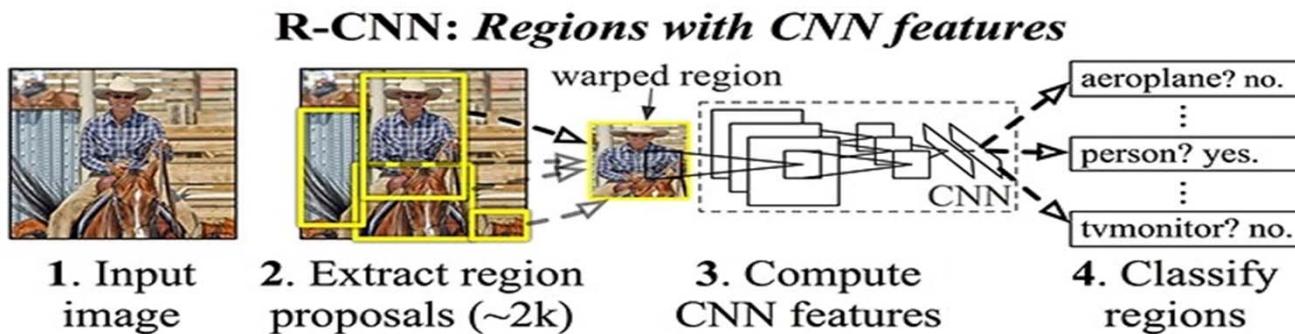
- Find “blobby” image regions that are likely to contain objects
- “Class-agnostic” object detector
- Look for “blob-like” regions



1 Feb 2016

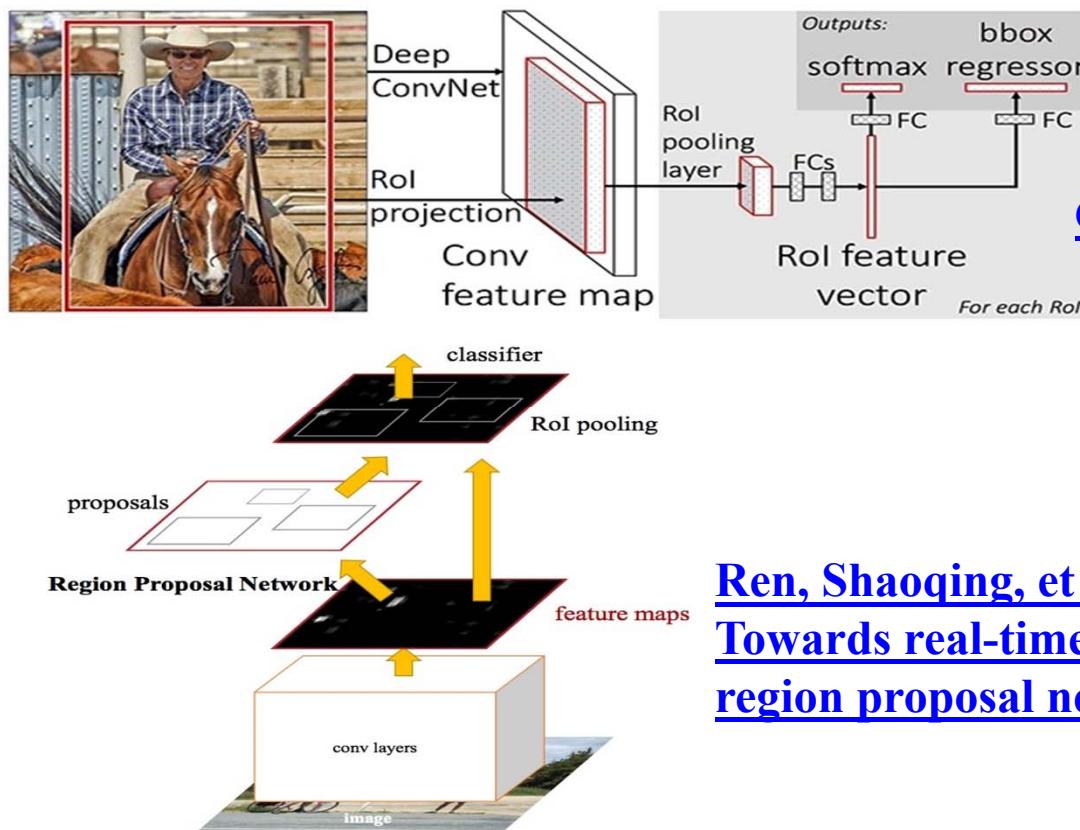
R-CNN (RP+CNN)

- The first CNN based detection model was R-CNN
- Combines: VGG and Region Proposals
- Boosted the object detection challenge by almost 50%.



[Girshick, Ross, et al. "Rich feature hierarchies for accurate object detection and semantic segmentation." 2014.](#)

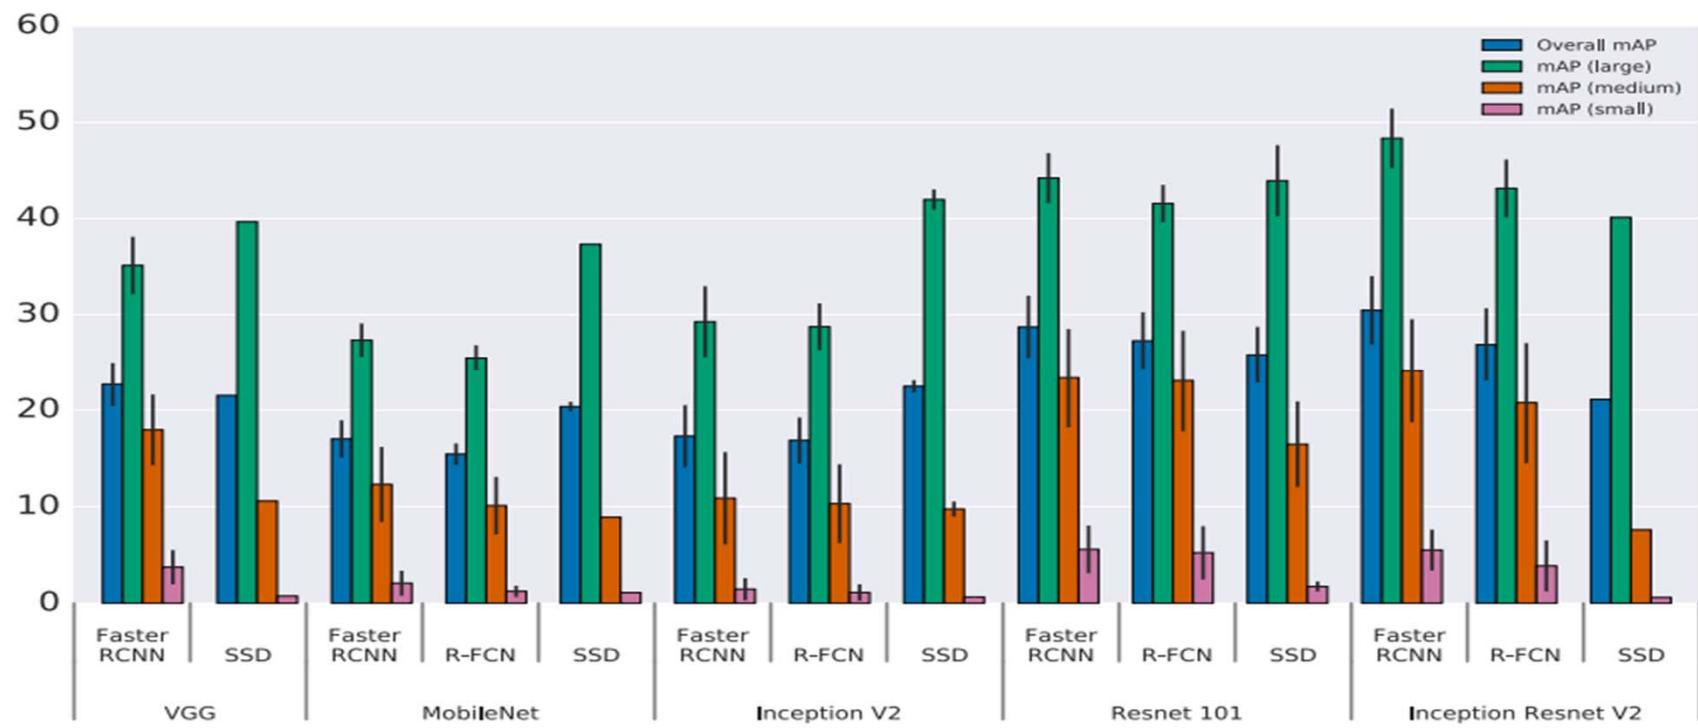
Fast R-CNN and Faster R-CNN



[**Girshick, Ross. "Fast R-CNN" 2015.**](#)

[**Ren, Shaoqing, et al. "Faster R-CNN: Towards real-time object detection with region proposal networks." 2015.**](#)

Faster R-CNN, SSD, R-FCN

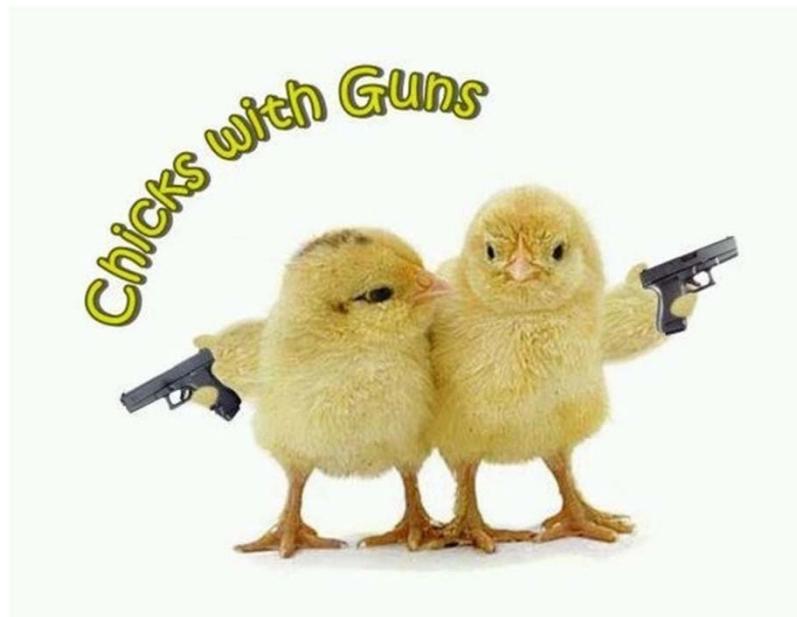


Outline

- What is detection?
- How detection models work?
- The state-of-the art detection models
- Case study 1
- Case study 2

Case study 3: Handgun alarm detection in videos

- Objective: Develop a fast and accurate pistol detection model in videos
- Methodology:
 - New database
 - Two models:
 - “VGG+ Sliding Window”
 - “VGG+ Region Proposal (Faster-RCNN)”

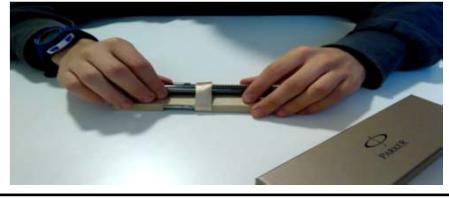


[Automatic Handgun Detection Alarm in Videos Using Deep Learning](#) R

Olmos, S Tabik, F Herrera. Neurocomputing, 275, 66-72

Case study 3: Handgun alarm detection in videos

A new database

Pistol class			
			
Negative class			
			

Results on the same dataset

Database	#TP	#FN	#TN	#FP	Precision	Recall	F1_measure
For Sliding window	97	207	298	6	94.17%	31.91%	47.67%
For Faster RCNN	304	0	247	57	81.21%	100.00%	91.43%

Results in images

Results on pictures



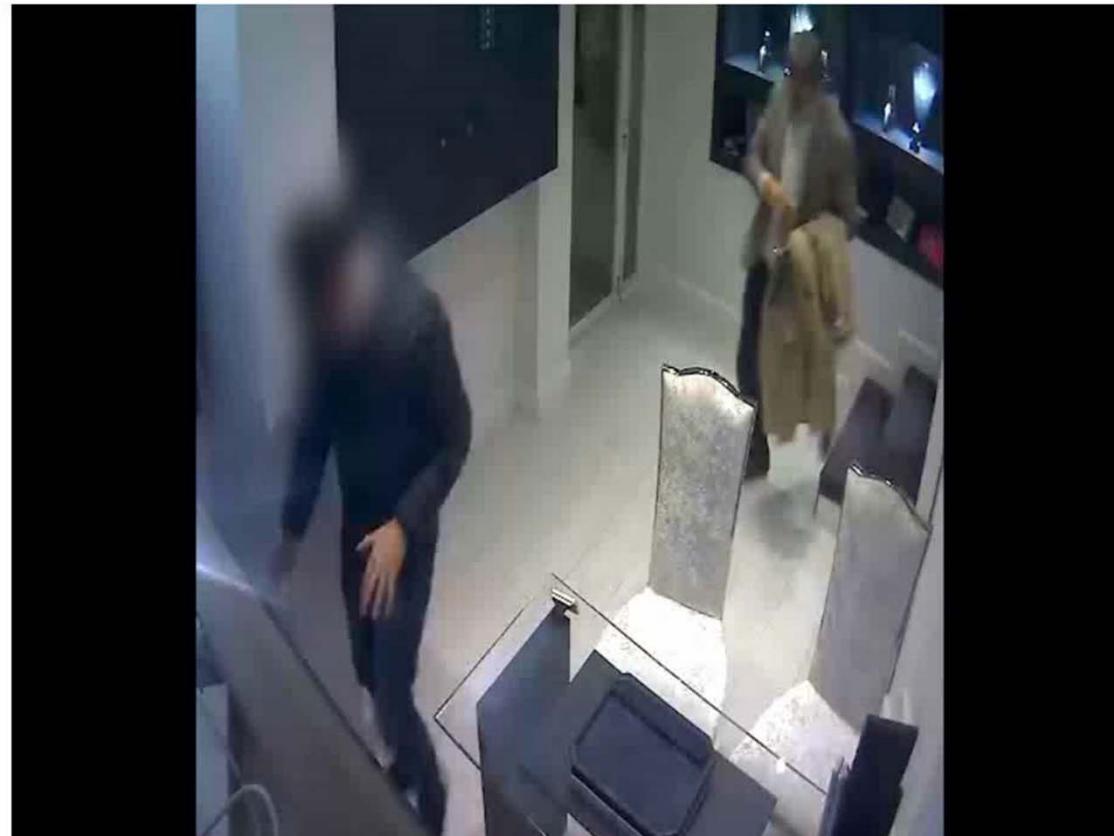
Results on youtube videos



Results on Skyfall



Results on real surveillance videos



Outline

- What is detection?
- How detection models work?
- The state-of-the art detection models
- Case study 1
- Case study 2

Case study 2: Protected plant species detection using Google-earth imagery

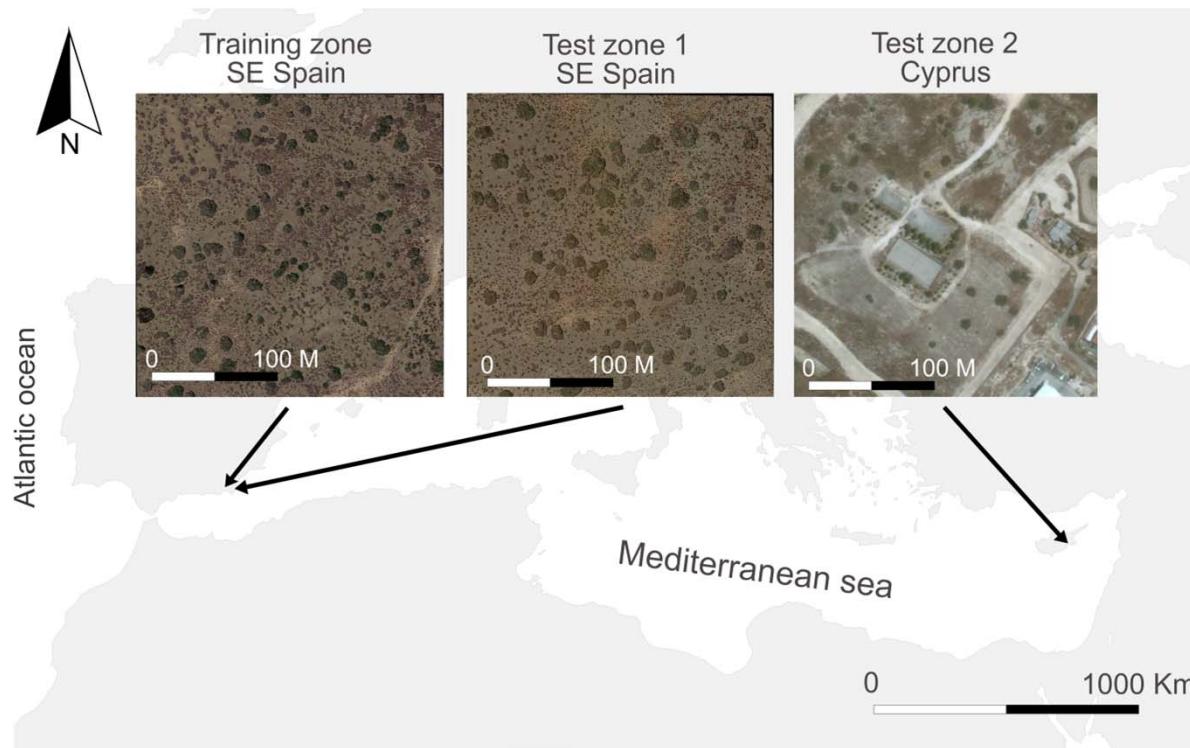
- Objective: *Ziziphus Lotus* detection in GE images
- Methodology:
 - Design new database to train CNNs
 - Three models: “CNNs+ Sliding Window” “CNNs+ Preprocessing” “OBIA”



[Deep-Learning Convolutional Neural Networks for scattered shrub detection with Google Earth Imagery](#) E Guirado, S Tabik, D Alcaraz-Segura, J Cabello, F Herrera. [Remote Sensing](#) 9(12): 1220 (2017)

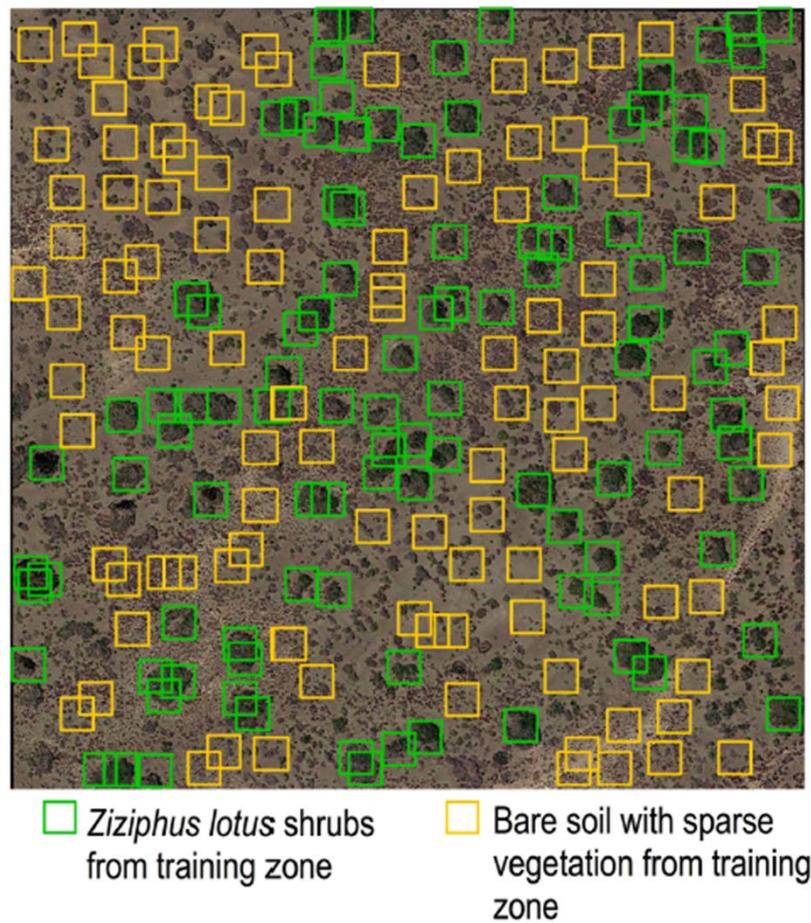
Case study 2: Protected plant species detection using Google-earth imagery

- Design new database to train CNNs

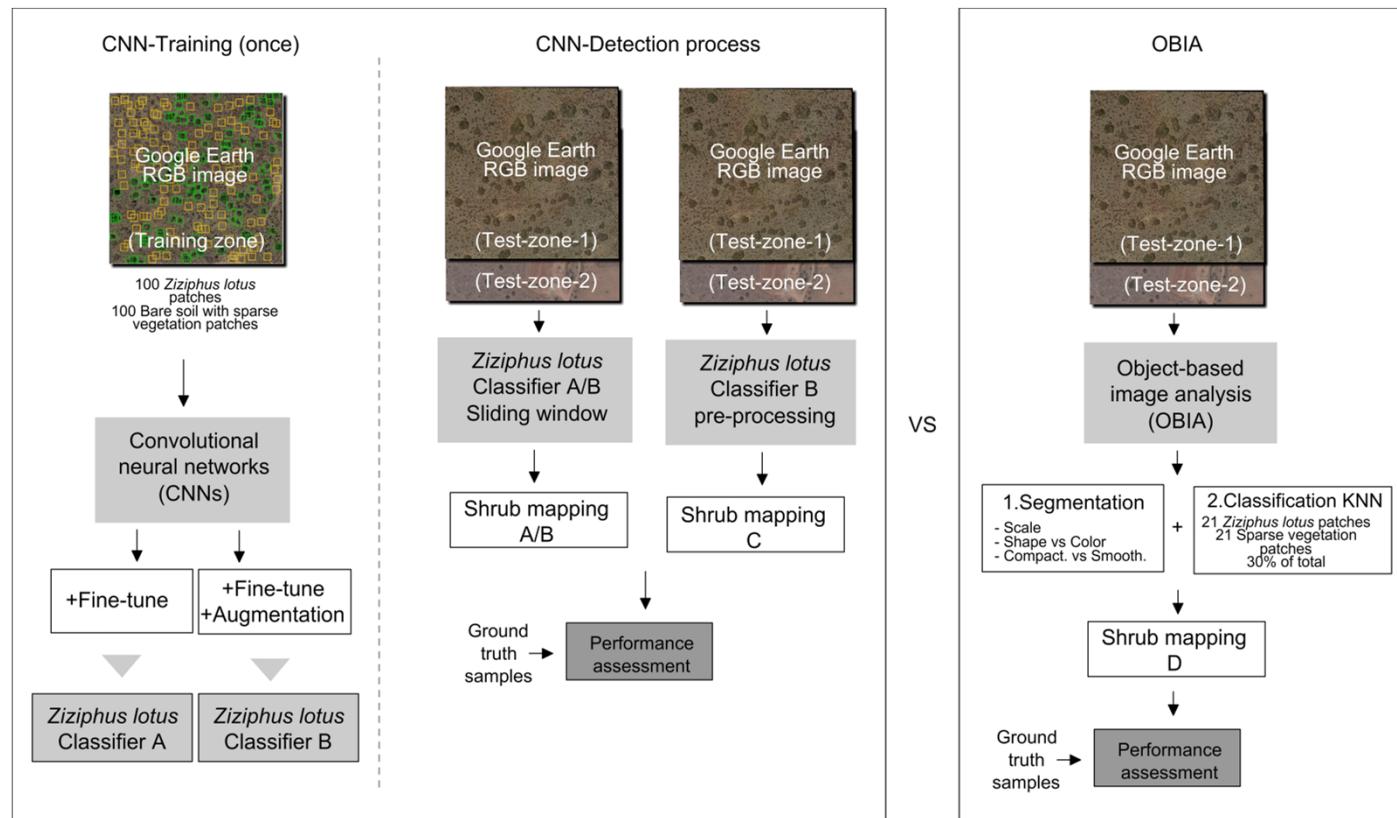


Case study 2: Protected plant species detection using Google-earth imagery

- Database for training the CNNs
- A two-class problem



Case study 2: Protected plant species detection using Google-earth imagery



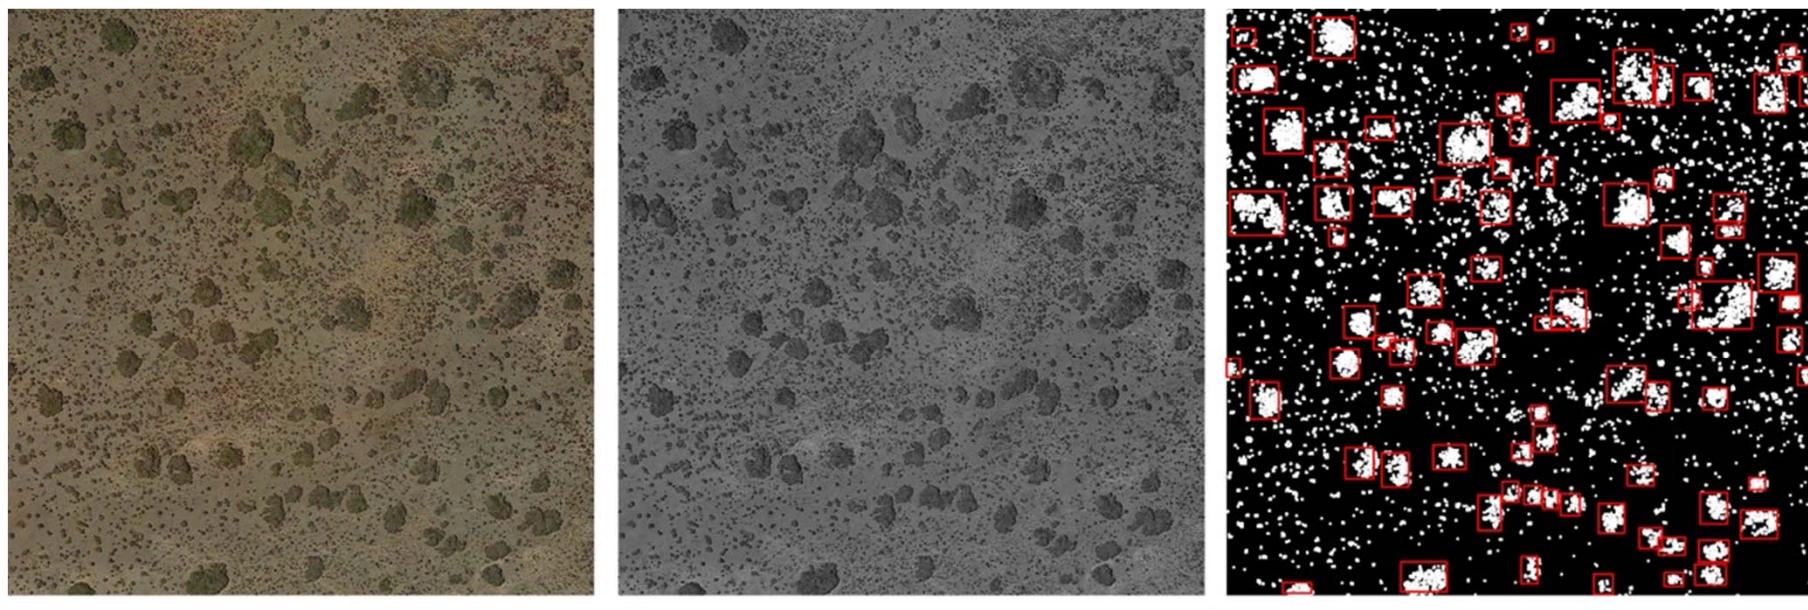
$$precision = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}},$$

$$recall = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}},$$

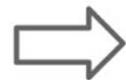
$$F1\ measure = 2 \times \frac{precision \times recall}{precision + recall}$$

Case study 2: Protected plant species detection using Google-earth imagery

Preprocessing:



1. RGB image to PAN image



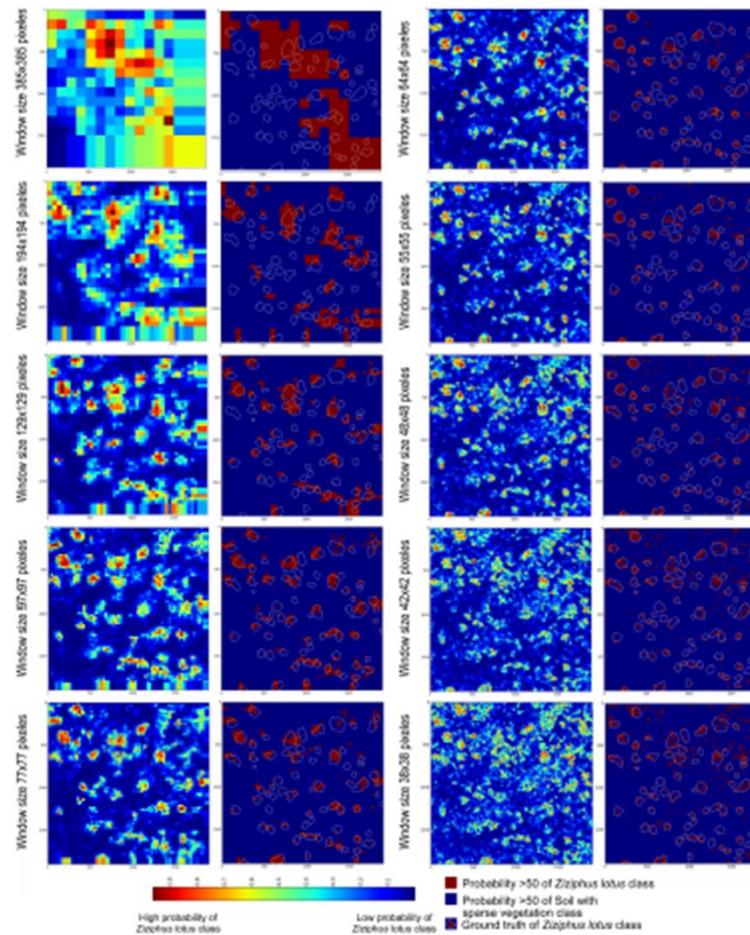
2. PAN image to gray threshold image



3. Gray threshold image to contour threshold candidate patches
(Pixel value < 100)

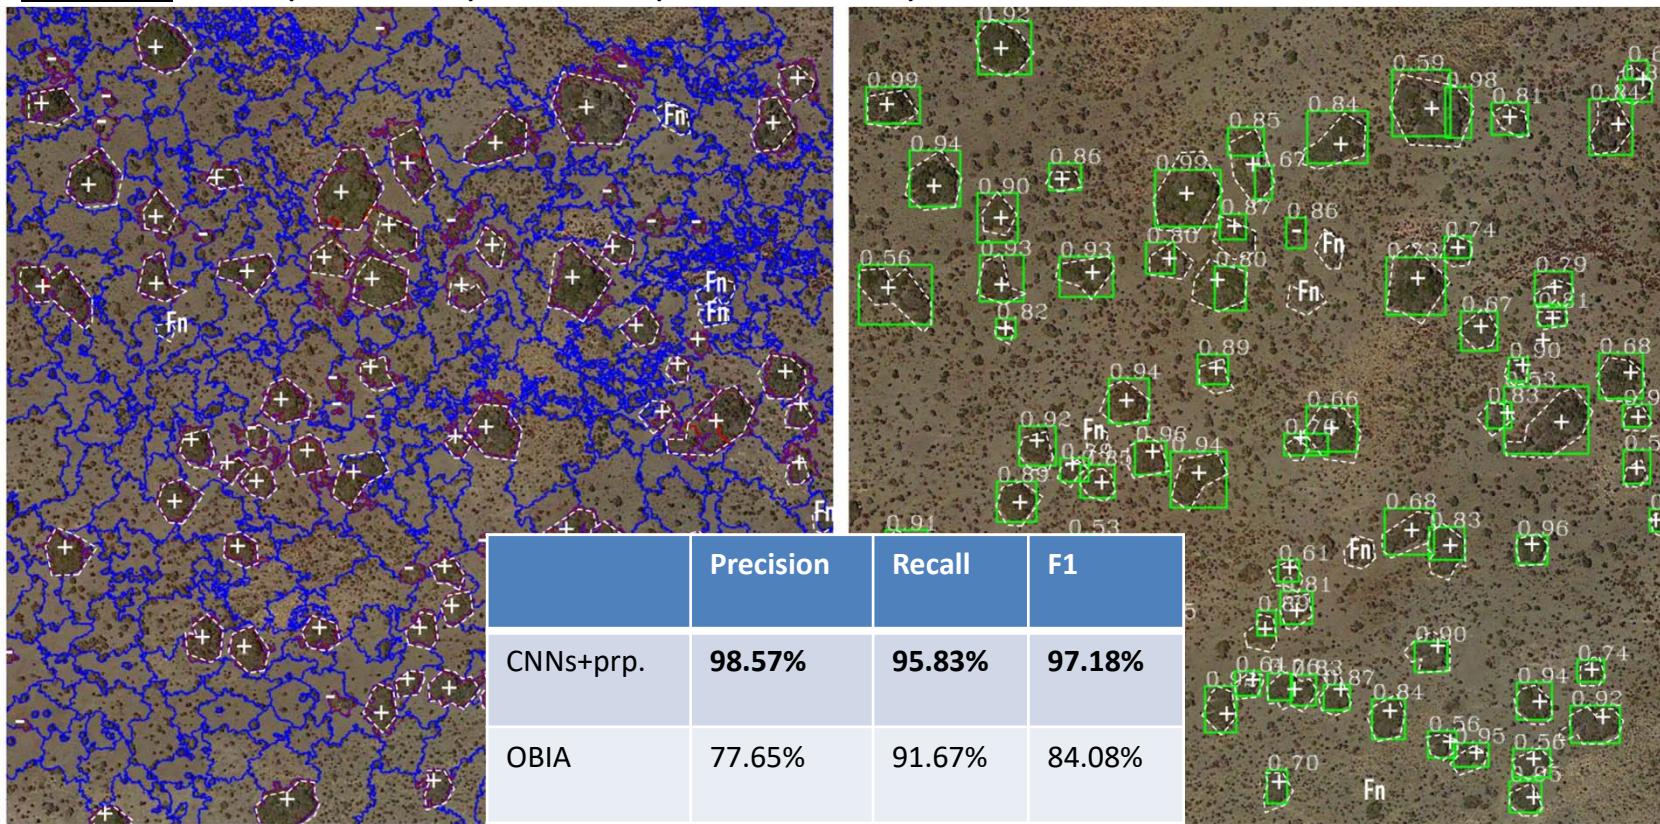
Case study 2: Protected plant species detection using Google-earth imagery

- Results
(GoogLeNet+sliding window)
- It takes: 291 min



Case study 2: Protected plant species detection using Google-earth imagery

Results: OBIA(12 hours) vs CNNs(35.4seconds)



Case study 2: Protected plant species detection using Google-earth imagery

Results: OBIA(12 hours) vs CNNs(24.1seconds)

